**omermahgoub** 1 year, 6 months ago

The correct answer is A: Add all users to a group. Grant the group the role of BigQuery user on the billing project and BigQuery dataViewer or projects that contain the data.

To make sure that no query costs are incurred on the projects that contain the data and allow users to query the datasets but not edit them, y should follow these steps:

Add all users to a group.
Grant the group the role of BigQuery user on the billing project. This will allow the group to run queries on BigQuery and incur costs on the bil project.
Grant the group the role of BigQuery dataViewer on the projects that contain the data. This will allow the group to view the datasets and run queries on them, but not edit them.

upvoted 1 times

**omermahgoub** 1 year, 6 months ago

The BigQuery Job User role (roles/bigquery.jobUser) and the BigQuery User role (roles/bigquery.user) have similar permissions, but they di in the scope of their permissions.

The BigQuery Job User role grants users the ability to create and modify query jobs, but it does not grant them the ability to run queries or incur costs on the project. This role is intended for users who need to create and manage query jobs, but who should not be able to run queries or incur costs.

The BigQuery User role grants users the ability to run queries and incur costs on the project, in addition to the ability to create and modify query jobs. This role is intended for users who need to run queries and incur costs on the project, as well as create and manage query jobs

upvoted 3 times

**omermahgoub** 1 year, 6 months ago

Here is a summary of the differences between the BigQuery Job User role and the BigQuery User role:

BigQuery Job User role (roles/bigquery.jobUser):

Can create and modify query jobs
Cannot run queries or incur costs on the project
BigQuery User role (roles/bigquery.user):

Can create and modify query jobs
Can run queries and incur costs on the project
If you want to grant users the ability to create and modify query jobs, but not run queries or incur costs on the project, you should use t BigQuery Job User role. If you want to grant users the ability to run queries and incur costs on the project, in addition to the ability to create and modify query jobs, you should use the BigQuery User role.

upvoted 2 times

**omermahgoub** 1 year, 6 months ago

Option B: Granting the group the roles of BigQuery dataViewer on the billing project and BigQuery user on the projects that contain the da will not allow the group to incur costs on the billing project and will not meet the requirements of the scenario.

Option C: Granting the group the roles of BigQuery jobUser on the billing project and BigQuery dataViewer on the projects that contain the data will not allow the group to incur costs on the billing project and will not meet the requirements of the scenario.

Option D: Granting the group the roles of BigQuery dataViewer on the billing project and BigQuery jobUser on the projects that contain the data will not allow the group to incur costs on the billing project and will not meet the requirements of the scenario.

upvoted 1 times

**Diwz** 2 months, 1 week ago

BigQuery User
(roles/bigquery.user)
When applied to a dataset, this role provides the ability to read the dataset's metadata and list tables in the dataset.

When applied to a project, this role also provides the ability to run jobs, including queries, within the project. A principal with this role can enumerate their own jobs, cancel their own jobs, and enumerate datasets within a project. Additionally, allows the creation of new datasets within the project; the creator is granted the BigQuery Data Owner role (roles/bigquery.dataOwner) on these new datasets.

Bigquery.user has potential to create a dataset inside the project and creates becomes owner of the dataset. This is not the requirement stated in the question scenario.
Answer is C

upvoted 1 times

**AzureDP900** 1 year, 8 months ago

C is right
Add all users to a group. Grant the group the roles of BigQuery jobUser on the billing project and BigQuery dataViewer on the projects that contain the data.

upvoted 1 times

**minmin2020** 1 year, 8 months ago

Selected Answer: C

C. Add all users to a group. Grant the group the roles of BigQuery jobUser on the billing project and BigQuery dataViewer on the projects that contain the data.

upvoted 1 times

**Vedjha** 1 year, 8 months ago

D is the answer:
Cloud BigQuery Roles
Cloud BigQuery IAM Roles
BigQuery Admin - bigquery.*
BigQuery Data Owner - bigquery.datasets.*, bigquery.models.*, bigquery.routines.*, bigquery.tables.* (Does NOT have access to Jobs!)
BigQuery Data Editor - bigquery.tables.(create/delete/export/get/getData/getIamPolicy/list/update/updateData/updateTag), bigquery.models.*, bigquery.routines.*, bigquery.datasets.(create/get/getIamPolicy/updateTag)
BigQuery Data Viewer - get/list bigquery.(datasets/models/routines/tables)
BigQuery Job User - bigquery.jobs.create
BigQuery User - BigQuery Data Viewer + get/list (jobs, capacityCommitments, reservations etc)
To see data, you need either BigQuery User or BigQuery Data Viewer roles
You CANNOT see data with BigQuery Job User roles
BigQuery Data Owner or Data Viewer roles do NOT have access to jobs!

upvoted 1 times

**kimharsh** 2 years, 4 months ago

C is the correct Answer ,
A is wrong because bq User Permission will allow you to edit the dataset, which is something that we don't want in this scenario.
B and D is wrong because "You want to make sure that no query costs are incurred on the projects that contain the data" so you don't want users to fire quires on the Project that contains the dataset , hence the "dataViewer" permission

https://cloud.google.com/bigquery/docs/access-control

upvoted 1 times

**victory108** 2 years, 5 months ago

C. Add all users to a group. Grant the group the roles of BigQuery jobUser on the billing project and BigQuery dataViewer on the projects that contain the data.

upvoted 4 times

**LoveT** 2 years, 5 months ago

C looks to be the correct answer

upvoted 2 times

**HenkH** 2 years, 5 months ago

Selected Answer: C

JobUser is the correct terminology for bq. Only read access to data sources is required.

upvoted 1 times

**HenkH** 2 years, 5 months ago

bq is using jobs - so "user" isn't specific enough, jobuser is.

upvoted 2 times

**elenamatay** 2 years, 5 months ago

Hence C

upvoted 1 times

Question #58 *Topic 1*

You have developed an application using Cloud ML Engine that recognizes famous paintings from uploaded images. You want to test the application and allow specific people to upload images for the next 24 hours. Not all users have a Google Account. How should you have users upload images?

A. Have users upload the images to Cloud Storage. Protect the bucket with a password that expires after 24 hours.

B. Have users upload the images to Cloud Storage using a signed URL that expires after 24 hours.

C. Create an App Engine web application where users can upload images. Configure App Engine to disable the application after 24 hours. Authenticate users via Cloud Identity.

D. Create an App Engine web application where users can upload images for the next 24 hours. Authenticate users via Cloud Identity.

**jcmoranp** Highly Voted 5 years, 1 month ago
Correct answer is B
upvoted 44 times

**tartar** 4 years, 4 months ago
B is ok
upvoted 8 times

**kumarp6** 4 years, 1 month ago

Signed URL ... B is correct

upvoted 3 times

**nitinz** 3 years, 9 months ago

B signed URL

upvoted 3 times

**MyPractice** Highly Voted 4 years, 11 months ago

Ans B
"When should you use a signed URL? In some scenarios, you might not want to require your users to have a Google account in order to acce Cloud Storage" "Signed URLs contain authentication information in their query string, allowing users without credentials to perform specific actions on a resource"
https://cloud.google.com/storage/docs/access-control/signed-urls

upvoted 27 times

**sim7243** Most Recent 1 month, 1 week ago

**Selected Answer: B**

Correct answer is B

upvoted 1 times

**gun123** 11 months, 2 weeks ago

**Selected Answer: B**

Correct answer is B

upvoted 1 times

**red_panda** 1 year, 5 months ago

**Selected Answer: B**

B is the answer

upvoted 1 times

**fussili** 1 year, 8 months ago

The correct answer is B.

A is not a good choice because it is not possible to set an expiration time for a password protected Cloud Storage bucket. This means that if user had the password, they would be able to upload images to the bucket even after the 24 hour period has expired.

B is the correct answer because a signed URL can be generated to allow specific users to upload images to Cloud Storage without requiring them to have a Google Account. The URL can be set to expire after 24 hours, which ensures that users can only upload images during the allowed time period.

C is not the best choice because it involves creating an App Engine web application, which is more complex than using Cloud Storage with a signed URL. Additionally, App Engine instances cannot be turned off programmatically, so it would not be possible to disable the application a 24 hours.

D option is similar to option C, but it involves creating an App Engine web application. This would add unnecessary complexity to the solution and it would not provide any additional benefits compared to using Cloud Storage with a signed URL.

upvoted 3 times

**omermahgoub** 1 year, 12 months ago

The correct answer is B: Have users upload the images to Cloud Storage using a signed URL that expires after 24 hours.

To allow specific users to upload images to Cloud Storage for testing your Cloud ML Engine application, and to not require all users to have a Google Account, you should use signed URLs. A signed URL is a URL that allows access to a specific resource in Cloud Storage, and that is valid for a specified period of time.

To create a signed URL that expires after 24 hours, you can use the gsutil signurl command. For example:

Copy code
gsutil signurl -d 24h service-account.json gs://bucket-name/object-name
This will generate a signed URL that allows users to upload an object to the specified bucket with the specified name, and that will only be va for 24 hours.

upvoted 3 times

**omermahgoub** 1 year, 12 months ago

Option A: Protecting the bucket with a password that expires after 24 hours would not be a secure or scalable solution, as it would require to distribute the password to all users and to update the password every 24 hours.

Option C: Creating an App Engine web application where users can upload images, and configuring App Engine to disable the application after 24 hours, would not allow users to upload images after the application is disabled.

Option D: Creating an App Engine web application where users can upload images for the next 24 hours and authenticating users via Clou Identity would not allow users to upload images if they do not have a Google Account.

upvoted 1 times

**megumin** 2 years, 1 month ago

**Selected Answer: B**

B is ok

upvoted 1 times

**AzureDP900** 2 years, 2 months ago

B is right, Signed URL's will help in this scnerio.

upvoted 1 times

**minmin2020** 2 years, 2 months ago

**Selected Answer: B**

B. Have users upload the images to Cloud Storage using a signed URL that expires after 24 hours.

upvoted 1 times

**mv2000** 2 years, 5 months ago

On 06/30/2022 Exam.

upvoted 2 times

**mygcpjourney2712** 2 years, 8 months ago

**Selected Answer: B**

signed url

upvoted 1 times

**vincy2202** 2 years, 11 months ago

B is the correct answer

upvoted 2 times

**haroldbenites** 3 years ago

Go for B.

upvoted 1 times

**MaxNRG** 3 years, 1 month ago

Question #59 *Topic 1*

Your web application must comply with the requirements of the European Union's General Data Protection Regulation (GDPR). You are responsible for the technical architecture of your web application. What should you do?

A. Ensure that your web application only uses native features and services of Google Cloud Platform, because Google already has various certifications and provides λ€pass-onλ€ compliance when you use native features.

B. Enable the relevant GDPR compliance setting within the GCPConsole for each of the services in use within your application.

C. Ensure that Cloud Security Scanner is part of your test planning strategy in order to pick up any compliance gaps.

D. Define a design for the security of data in your web application that meets GDPR requirements.

**AWS56** Highly Voted 3 years, 5 months ago

Agree D

upvoted 17 times

**AshokC** Highly Voted 2 years, 9 months ago

D - https://cloud.google.com/security/gdpr
The GDPR lays out specific requirements for businesses and organizations who are established in Europe or who serve users in Europe. It:

Regulates how businesses can collect, use, and store personal data
Builds upon current documentation and reporting requirements to increase accountability
Authorizes fines on businesses who fail to meet its requirements

upvoted 15 times

**omermahgoub** Most Recent 6 months ago

The correct answer is option D: Define a design for the security of data in your web application that meets GDPR requirements.

The General Data Protection Regulation (GDPR) is a comprehensive data protection law that applies to any company that processes the pers data of individuals in the European Union (EU). As the technical architect of your web application, it is your responsibility to ensure that the application is compliant with GDPR requirements.

upvoted 4 times

**omermahgoub** 6 months ago

Option A: While it is true that Google has various certifications and provides pass-on compliance when you use native features, simply usi native features and services of Google Cloud Platform is not sufficient to ensure compliance with GDPR. You still need to implement appropriate controls and safeguards to protect personal data and meet GDPR requirements.

Option B: Enabling the relevant GDPR compliance setting within the GCP console for each of the services in use within your application m help ensure compliance with GDPR, but it is not sufficient on its own. You still need to implement appropriate controls and safeguards to protect personal data and meet GDPR requirements.

Option C: Using Cloud Security Scanner as part of your test planning strategy can help identify potential security vulnerabilities and compliance gaps in your web application, but it is not sufficient on its own to ensure compliance with GDPR. You still need to implement appropriate controls and safeguards to protect personal data and meet GDPR requirements.

upvoted 8 times

**megumin** 7 months, 2 weeks ago

Selected Answer: D

D is ok

upvoted 1 times

**AzureDP900** 8 months ago

Define a design for the security of data in your web application that meets GDPR requirements. D is right

upvoted 1 times

**[Removed]** 1 year, 4 months ago

Selected Answer: D

I got similar question on my exam.

upvoted 6 times

**vincy2202** 1 year, 5 months ago

D is the correct answer

upvoted 1 times

**haroldbenites** 1 year, 6 months ago

Go for D

upvoted 1 times

**joe2211** 1 year, 6 months ago

**Selected Answer: D**

vote D

upvoted 1 times

**MaxNRG** 1 year, 7 months ago

D – Define a design for the security of data in your web app that meets GDPR requirements.

upvoted 1 times

**MikeB19** 1 year, 9 months ago

A is wrong D is correct. The q refers is “Microsoft sql” not “MySQL”. App replication in MSsql is achieved with Availability Groups within MSs
https://docs.microsoft.com/en-us/sql/database-engine/availability-groups/windows/overview-of-always-on-availability-groups-sql-server?view=sql-server-ver15

upvoted 1 times

**MamthaSJ** 1 year, 11 months ago

Answer is D

upvoted 3 times

**victory108** 2 years, 1 month ago

D. Define a design for the security of data in your web application that meets GDPR requirements.

upvoted 2 times

**un** 2 years, 1 month ago

D is correct

upvoted 1 times

**Ausias18** 2 years, 2 months ago

Answer is D

upvoted 1 times

**lynx256** 2 years, 2 months ago

D is ok

upvoted 1 times

**CloudGenious** 2 years, 4 months ago

you should design your app such that they meet GDPR req. As a customer google cloud, GDPR should be part of protection strategy .So the is D.

upvoted 3 times

Question #60 *Topic 1*

You need to set up Microsoft SQL Server on GCP. Management requires that there's no downtime in case of a data center outage in any of the zones within a

zones within a
GCP region. What should you do?

A. Configure a Cloud SQL instance with high availability enabled.

B. Configure a Cloud Spanner instance with a regional instance configuration.

C. Set up SQL Server on Compute Engine, using Always On Availability Groups using Windows Failover Clustering. Place nodes in different subnets.

D. Set up SQL Server Always On Availability Groups using Windows Failover Clustering. Place nodes in different zones.

**learningpv** Highly Voted 4 years, 10 months ago
A seems correct.
"... high availability (HA) configuration for Cloud SQL instances... A Cloud SQL instance configured for HA is also called a regional instance an located in a primary and secondary zone within the configured region.
In the event of an instance or zone failure, this configuration reduces downtime, and your data continues to be available to client applications.
upvoted 61 times

**mrealtor** 2 years, 8 months ago
You need to set up a Microsoft SQL server. Why are we talking about Cloud SQL
upvoted 5 times

**tycho** 2 years, 4 months ago
and what is Cloud SQL -> a managed service for MySQL, Posrgers, and MS SQL server
upvoted 10 times

**Ric350** 4 months ago
Actually, although Cloud SQL offers high availability configurations, it currently doesn't support Microsoft SQL Server as one of its managed database engines. It primarily focuses on MySQL, PostgreSQL, and SQL Server (but not the full Microsoft SQL Server). An the question clearly states "you need to set up Microsoft SQL." Very tricky question. The answer is D
upvoted 3 times

**diluviouniv** 3 years, 5 months ago
but it says: you need to setup SQL Server
upvoted 12 times

**learningpv** 4 years, 10 months ago
It applies for MySQL and HA is not available for MS SQL
upvoted 5 times

**cetanx** 4 years, 5 months ago
It is available, please see;
https://cloud.google.com/sql/docs/sqlserver/high-availability?_ga=2.30855355.-503483612.1582800507
Also a video from Google;
https://youtu.be/vMUpNoukwnM
upvoted 11 times

**Jos** 4 years, 10 months ago
Yes it is available, its in beta, but when creating a "SQL Server 2017 Standard" in Cloud SQL menu you can chose single one or HA (regional).
upvoted 3 times

**tocsa** 6 months, 2 weeks ago
The problem is that these questions are ancient (talking about StackDriver all the time for example, it was rebranded in 2020!!!). So unfortunately we need to think of "What did the professor think 2-4 years ago" when this question was created. Otherwise I'd say A the best all day!
upvoted 1 times

**AmitAr** 2 years, 7 months ago
D is correct.
Question is - "no downtime while installing MS SQL" , not on choosing or replacing with GCP product. I agree A is good solution for this requirements.. however question is not on choosing database.. its for HA.. so I will choose D.
upvoted 4 times

**SMS** Highly Voted 4 years, 9 months ago

Answer is A. Cloud SQL supports SQL Server and selecting high availability provides automatic failover within a region.

upvoted 30 times

**plumbig11** Most Recent 5 days, 14 hours ago

**Selected Answer: D**

thinking about failover D is a better option

upvoted 1 times

**Srrb20** 1 week, 3 days ago

**Selected Answer: D**

Cloud SQL with HA is not sufficient because it does not provide the level of control and multi-zone redundancy that Always On Availability Groups on Compute Engine can deliver. Always On is the recommended solution for enterprise-grade high availability and disaster recovery ir SQL Server deployments.

upvoted 1 times

**klayytech** 2 weeks, 1 day ago

**Selected Answer: A**

Cloud SQL offers high availability configurations, it currently support Microsoft SQL Server
please see;
https://cloud.google.com/sql/docs/sqlserver/high-availability?_ga=2.30855355.-503483612.1582800507
Also a video from Google;
https://youtu.be/vMUpNoukwnM

upvoted 1 times

**valgorodetsky** 3 weeks, 5 days ago

**Selected Answer: D**

- HA - has minimal downtime
- windows-server-failover-clustering has zero, but much harder to setup: https://cloud.google.com/compute/docs/tutorials/running-windows-server-failover-clustering

upvoted 2 times

**ngeorgiev2** 1 month, 1 week ago

A looks more relevant - Multiple zones (Highly available)
Automatic failover to another zone within your selected region. Recommended for production instances. Increases cost.

upvoted 1 times

**SerGCP** 1 month, 1 week ago

**Selected Answer: D**

Cloud SQL with HA is good and may reduce downtime. But the requirement is "no dowtime" so D.

upvoted 2 times

**sim7243** 1 month, 1 week ago

A is the option,

upvoted 1 times

**bd311b9** 1 month, 2 weeks ago

D is the right answer because high availability is created by putting a failover instance in a different zone.

upvoted 1 times

**nareshthumma** 1 month, 3 weeks ago

The correct approach is: D

Set up SQL Server Always On Availability Groups using Windows Failover Clustering. Place nodes in different zones.

Here's why this is the best option:

• SQL Server Always On Availability Groups: This solution provides high availability by automatically failing over to another node in the event c failure. It's specifically designed for SQL Server and ensures minimal downtime in case of outages.

• Windows Failover Clustering: By configuring Windows Failover Clustering with Always On Availability Groups, you can achieve high availabil by ensuring that the SQL Server can failover to another node in case of a zone or node failure.

• Placing nodes in different zones: By deploying nodes in different zones within the same region, you ensure that your setup is protected from potential zone-level outages. If one zone experiences a failure, the other zone can take over without downtime.

upvoted 1 times

**selected** 1 month, 3 weeks ago

use google docs instead of LLMs

upvoted 1 times

**JohnJamesB1212** 3 months, 1 week ago

**Selected Answer: D**

The correct answer is D. Set up SQL Server Always On Availability Groups using Windows Failover Clustering. Place nodes in different zones.

Here's why:

SQL Server Always On Availability Groups is a high-availability and disaster recovery solution for SQL Server that works across multiple zones ensuring minimal downtime in case of a data center outage within a region.
By placing the nodes in different zones, you ensure that the database remains accessible even if one zone goes down, meeting the requireme of no downtime in case of a zone failure.

upvoted 4 times

**JohnJamesB1212** 3 months, 1 week ago

The other options are less suitable for this scenario:

A. Cloud SQL with high availability enabled would not work because Cloud SQL supports MySQL, PostgreSQL, and SQL Server but has limitations for high-availability setups with SQL Server compared to Always On Availability Groups.
B. Cloud Spanner is a Google-native distributed database solution and not directly related to SQL Server.
C mentions using different subnets, but you specifically need to place the nodes in different zones to ensure availability across multiple da centers within a region.
Thus, D is the most accurate solution for ensuring no downtime in the event of a zone failure while using Microsoft SQL Server on GCP.

upvoted 2 times

**desertlotus1211** 4 months, 2 weeks ago

Cloud SQL for SQL Server is a fully managed relational database service from Google Cloud that allows users to set up, maintain, and manag Microsoft SQL Server databases in the cloud. Cloud SQL also supports MySQL and PostgreSQL

upvoted 2 times

**desertlotus1211** 4 months, 2 weeks ago

Answer A

upvoted 1 times

**nhatne** 5 months, 3 weeks ago

**Selected Answer: A**

go for A, Why not use the available option?

upvoted 1 times

**ccpmad** 6 months, 1 week ago

**Selected Answer: A**

It is A,
Who say D, better not to work in cloud, but yes in on premises windows servers LOL

upvoted 2 times

**hitmax87** 7 months ago

**Selected Answer: A**

A because of requirement: "no downtime in case of a data center outage in any of the zones within a GCP region", Spanner is for multi-region presence.

upvoted 1 times

**MarkosMuche** 8 months ago

Definitely A. Google always favors managed services over services you manage. It is a lot of work to set up and maintain your own server.

upvoted 3 times

Question #61 *Topic 1*

The development team has provided you with a Kubernetes Deployment file. You have no infrastructure yet and need to deploy the application. What should you do?

A. Use gcloud to create a Kubernetes cluster. Use Deployment Manager to create the deployment.

B. Use gcloud to create a Kubernetes cluster. Use kubectl to create the deployment.

C. Use kubectl to create a Kubernetes cluster. Use Deployment Manager to create the deployment.

D. Use kubectl to create a Kubernetes cluster. Use kubectl to create the deployment.

**MeasService** Highly Voted 4 years, 8 months ago

It has to be B. gcloud for creating cluster and kubectl for creating deployment

upvoted 54 times

**KouShikyou** Highly Voted 4 years, 8 months ago

May I ask why C is correct?
I thought B was correct.

upvoted 26 times

**nitinz** 3 years, 3 months ago

B, gcloud to manage GKE and to manage pods use kubctl.

upvoted 2 times

**kumarp6** 3 years, 7 months ago

B is correct, when you create a nodes in GKE you use gcloud rather than kubectl...

upvoted 4 times

**res3** 3 years, 11 months ago

yeap, gcloud command to create K8s cluster https://cloud.google.com/kubernetes-engine/docs/how-to/creating-a-cluster

upvoted 4 times

**tartar** 3 years, 10 months ago

B is ok

upvoted 9 times

**simiramis221** Most Recent 6 months ago

Answer is B %100

upvoted 1 times

**vamgcp** 1 year, 4 months ago

Create a Google Kubernetes Engine (GKE) cluster: You can use the Google Cloud Console or the gcloud command-line tool to create a GKE cluster, which will provide the underlying infrastructure for running your application.

Deploy the application to the cluster: You can use the kubectl command-line tool to apply the Kubernetes Deployment file provided by the development team to the cluster.kubectl apply -f deployment.yaml

upvoted 3 times

**megumin** 1 year, 7 months ago

Selected Answer: B

B is ok

upvoted 1 times

**Mahmoud_E** 1 year, 8 months ago

Selected Answer: B

is the correct answer https://cloud.google.com/kubernetes-engine/docs/how-to/creating-a-cluster

upvoted 1 times

**Mahmoud_E** 1 year, 8 months ago

B is the correct answer cluster https://cloud.google.com/kubernetes-engine/docs/how-to/creating-a-cluster

upvoted 1 times

**AzureDP900** 1 year, 8 months ago

B is right

upvoted 1 times

**SAMBIT** 2 years, 3 months ago

Kubctle comes live only when cluster has been created in the cloud console using cloud command

upvoted 2 times

**ghadxx** 2 years, 4 months ago

Selected Answer: B

Deployment Manager is used to automate the process of provisioning infrastructure. Therefore, gcloud and Deployment Manager do the same thing. Meanwhile, kubectl is used to run commands against an already created cluster.

upvoted 9 times

**haroldbenites** 2 years, 6 months ago

Go for B.
gcloud for create clusters.
kubectl is used when the cluster already has been created. For example to create deployments.
Kubectl has configured a config file where is specified the default cluster.

upvoted 2 times

**vincy2202** 2 years, 6 months ago

B is correct

upvoted 1 times

**Zinhle** 2 years, 7 months ago

Hi all may someone please share the link for the bank of questions because I cannot seem to locate them.

thank you

upvoted 1 times

**MaxNRG** 2 years, 7 months ago

B – use gcloud to create cluster, use kubectl to create a deployment.
https://cloud.google.com/kubernetes-engine/docs/how-to/creating-a-zonal-cluster
In fact, kubectl run creates a deployment.
https://cloud.google.com/kubernetes-engine/docs/tutorials/hello-app

upvoted 2 times

**ale183** 2 years, 9 months ago

Question for all , do we know if only new questions are part of the bank for new exam? Have any of the old questions appeared on new exam

upvoted 3 times

**xaliq** 2 years, 9 months ago
B is corrent
upvoted 1 times

**Raja101** 2 years, 9 months ago
Why not A ?
upvoted 3 times

Question #62 *Topic 1*

You need to evaluate your team readiness for a new GCP project. You must perform the evaluation and create a skills gap plan which incorporates the business goal of cost optimization. Your team has deployed two GCP projects successfully to date. What should you do?

A. Allocate budget for team training. Set a deadline for the new GCP project.

B. Allocate budget for team training. Create a roadmap for your team to achieve Google Cloud certification based on job role.

C. Allocate budget to hire skilled external consultants. Set a deadline for the new GCP project.

D. Allocate budget to hire skilled external consultants. Create a roadmap for your team to achieve Google Cloud certification based on job role.

**KouShikyou** Highly Voted 5 years, 1 month ago
B is correct.
upvoted 46 times

**kumarp6** 4 years, 1 month ago
Yes it is
upvoted 1 times

**passnow** 5 years ago
I would agree with you because the question says create a skills gap plan
upvoted 4 times

**nitinz** 3 years, 9 months ago
B, looks like cooked up question. Not gonna show up on actual test. Even if it does show up, its not market.
upvoted 10 times

**ACE_ASPIRE** 3 years, 2 months ago
exactly
upvoted 1 times

**mawsman** Highly Voted 4 years, 10 months ago
I think it's B. "You must perform the evaluation and create a skills gap plan incorporates the business goal of cost optimization." The goal is to cost optimize - they might have deployed 2 projects but are they cost optimized? I think the only way to evaluate the skills gap in cost optimization is to make them get certified and use the results to determine cost optimization skills gap. Quickly pushing another project deadl would not help with cost optimization.
upvoted 21 times

**Smart** 4 years, 10 months ago
Agreed. How is setting up a GCP project deadline helping towards skill gap and cost optimization.
upvoted 3 times

**maxdanny** Most Recent 3 months, 2 weeks ago

**Selected Answer: B**

Explanation:

Training and certification based on specific job roles (e.g., Cloud Architect, Data Engineer) will ensure your team has the necessary skills for th new GCP project and will help align their capabilities with cost optimization strategies.
Since your team has already successfully deployed two GCP projects, upskilling them with targeted training is a more cost-effective and sustainable solution than relying on external consultants.
This option balances both team readiness and the business goal of cost optimization, while building long-term internal expertise.

upvoted 2 times

**joecloud12** 4 months, 2 weeks ago

**Selected Answer: B**

b is obvious..u need to address the skill gap

upvoted 1 times

**geekywitcher** 5 months, 2 weeks ago

**Selected Answer: B**

B is correct.

upvoted 1 times

**Robert0** 6 months, 3 weeks ago

**Selected Answer: B**

B makes the most sense. If readiness is the goal, it makes sense to invest in formation

upvoted 1 times

**GianpiGale** 8 months ago

C and D excluded because the teams have successfully deployed 2 GCP project to date, seems best option is to train them

upvoted 1 times

**Teckexam** 11 months ago

**Selected Answer: B**

B: this is logical answer. Setting project deadline or getting external consultants is not a good planning option.

upvoted 1 times

**JPA210** 1 year, 2 months ago

Answer A is more realistic. Certification is important, but not to be in a question in the exam.

upvoted 1 times

**AdityaGupta** 1 year, 2 months ago

**Selected Answer: B**

You need to evaluate your team readiness for a new GCP project. You must perform the evaluation and create a skills gap plan which incorporates the business goal of cost optimization.

The business Goal (Long Term Goal) is cost optimization, hence investment will be needed. Option A says setting up deadlines, it will not help cost optimization. But giving training and obtaining certification will do.

upvoted 2 times

**piiizu** 1 year, 2 months ago

In the short run, A seems like the best option because this would enable a quicker transition to the cloud which is the business Goal of cost optimization. In terms of cost, certification will cost the organisation more.

In the long run, B is the right option I would recommend as an architect. This would reduce the skill gap, increase proficiency and ensure repeatability. The organisation will incur more but will reap the reward... Except they have a bad culture and the staff resign right after getting t certification. Lol.

upvoted 2 times

**Frusci** 1 year, 3 months ago

**Selected Answer: B**

B, better to train your team than hire, and setting a deadline would defeat the purpose of evaluating the team's readiness.

upvoted 2 times

**PKookNN** 1 year, 6 months ago

**Selected Answer: B**

B is a better answer

upvoted 2 times

**kaaiden** 1 year, 7 months ago

**Selected Answer: A**

GCP partnership need 3 project

upvoted 1 times

**TheCloudGuruu** 1 year, 7 months ago

**Selected Answer: B**

B is the right answer

upvoted 2 times

**grejao** 1 year, 8 months ago

It shows more like a trick than a really question.
Furthermore are this question relevant for certification?

upvoted 1 times

**telp** 1 year, 9 months ago

**Selected Answer: B**

You are in a google exam. Always choose certification for your teams.

upvoted 3 times

Question #63 *Topic 1*

You are designing an application for use only during business hours. For the minimum viable product release, you'd like to use a managed product that automatically `scales to zero` so you don't incur costs when there is no activity.
Which primary compute resource should you choose?

A. Cloud Functions

B. Compute Engine

C. Google Kubernetes Engine

D. AppEngine flexible environment

**abirroy** Highly Voted 2 years, 3 months ago

**Selected Answer: A**

A. Cloud Functions - managed service scales down to 0
B. Compute Engine - not a managed service
C. Google Kubernetes Engine - not a managed service and wont scale down to 0
D. AppEngine flexible environment - managed service but wont scale down to 0

upvoted 25 times

**AzureDP900** 2 years, 2 months ago

Agree with A

upvoted 3 times

**NiveusSol** 9 months, 3 weeks ago

GKE is a managed service.

upvoted 4 times

**victory108** Highly Voted 2 years, 11 months ago

A. Cloud Functions

upvoted 11 times

**vpatiltech** 2 years, 10 months ago

Cloud function is more for event driven computing. We surely need k8s or app engine. Flex always have 1 instance running. So GKE shoul the option

upvoted 5 times

**YAS007** 2 years, 7 months ago

from the doc :https://cloud.google.com/kubernetes-engine/docs/concepts/cluster-autoscaler

Note: If you specify a minimum of zero nodes, an idle node pool can scale down completely. However, at least one node must always b available in the cluster to run system Pods.

upvoted 1 times

**6721sora** 2 years, 3 months ago

But no cost for System/Control nodes

upvoted 1 times

**Toothpick** Most Recent 4 months, 3 weeks ago

The only correct answer here is cloud run which isn't listed as an option.
(A) can scale to zero, but is mean't for event driven workloads,not full sized applications, especially in the routing department
(B) Needs manual stoppage, managed groups also keep one instance alive always
(C) You're billed for the control plane even if the node pools are empty
(D) Flex can't scale down to zero currently

upvoted 3 times

**MiguelMiguel** 1 year, 1 month ago

This question is confused. It's seems that the answer is cloud function because is the only platform that can scale 0 natively. If you want to sc to zero k8s, you have to create a solution based in scheduler, function and pub sub. Compute Engine is not a managed service, and app engi flex doesn't scale to 0. I would go to A but it's not clear.

upvoted 1 times

**LaxmanTiwari** 1 year, 7 months ago

agree A

upvoted 1 times

**grejao** 1 year, 8 months ago

Selected answer: D
I choosed D, it appears that we do not. have a right answer here.
A. Cloud Functions - its more for event driven computing, not for full application
B. Compute Engine - not a managed service
C. Google Kubernetes Engine - not a managed service and wont scale down to 0
D. AppEngine flexible environment - Only Standard App Engine can scale to 0.

upvoted 3 times

**Bedmed** 1 year, 7 months ago

yes, but only Standard environment, not flexible environment

upvoted 3 times

**GCPAnji** 1 year, 8 months ago

For an application that is only used during business hours and needs to scale to zero during periods of inactivity to minimize costs, a good choice would be a Function-as-a-Service (FaaS) product like AWS Lambda or Google Cloud Functions.

upvoted 1 times

**jlambdan** 1 year, 9 months ago

**Selected Answer: B**

A a cloud function is not an application

B compute engine via MIG you can use an autoscaler with a schedule.
https://cloud.google.com/compute/docs/autoscaler/scaling-schedules
You then can go from 0 to more instance when required

C K8s is two complex for this.
You can have an autoscaler for the cluster in order to get the node number to 0, but it require the node to have no pods running. So you have configure your deployments and all your workload to scale to 0 too.
Other interference will be pod affinity, anti-affinity, disruption budget or unmanaged pod preventing pod eviction from node. But if the pod is n evicted, the node cannot be deleted.
"autoscaler respects scheduling and eviction rules set on Pods. These restrictions can prevent a node from being deleted by the autoscaler. "
https://cloud.google.com/kubernetes-engine/docs/concepts/cluster-autoscaler#scheduling-and-disruption

D flexible cannot scale to 0

upvoted 1 times

**jrisl1991** 1 year, 2 months ago

App Engine is not an application on its own either, nor B or C options. B and C are not managed products; "managed" means that all of th infrastructure work (such as setting up an autoscheduler or autoscaler) will be managed by Google and not by the user. Using an IaaS solu for a managed use case is already contradictory.

App Engine Flex will always have at least 1 instance running.

While Cloud Functions is not an application itself, it's the only resource that can scale to 0. Plus, the requirement states that they will have application running during business hours, but they never mentioned that the app will only be running in GCP or in Cloud Functions. They could be running their app anywhere else, and only call Cloud Functions when needed.

upvoted 1 times

**theBestStudent** 1 year ago

That is not the concept of managed: https://cloud.google.com/blog/topics/developers-practitioners/serverless-vs-fully-managed-whats difference

upvoted 1 times

**CGS22** 1 year, 9 months ago

**Selected Answer: A**

A. Cloud Functions

upvoted 2 times

**WFCheong** 1 year, 11 months ago

Why not D? App Engines also can scale down to zero when there is no activity. https://cloud.google.com/appengine/docs/the-appengine-environments#:~:text=Intended%20to%20run%20for%20free,when%20there%20is%20no%20traffic. Intended to run for free or at very low cost, where you pay only for what you need and when you need it. For example, your application can scale to 0 instances when there is no traffic.

upvoted 3 times

**CkWongCk** 1 year, 10 months ago

There are 2 mode for App engines, standard and flexible.

The standard environment can scale from zero instances up to thousands very quickly. In contrast, the flexible environment must have at le one instance running for each active version and can take longer to scale out in response to traffic.

upvoted 4 times

**thamaster** 1 year, 12 months ago

**Selected Answer: A**

the only to scale to 0 is A

upvoted 2 times

**omermahgoub** 1 year, 12 months ago
The correct answer is A. Cloud Functions.

Cloud Functions is a serverless compute service that lets you run code without provisioning or managing infrastructure. One of the key benefi of using Cloud Functions is that it automatically scales to meet the demands of your workload and automatically scales down to zero when th is no activity. This means that you only pay for the compute resources that you consume, which can help to reduce costs when your applicati is not in use. Additionally, Cloud Functions is easy to use and allows you to deploy your code with minimal effort, making it a good choice for minimum viable product release.
upvoted 3 times

**backhand** 2 years, 4 months ago
vote A
this is easy one. key word: managed product, scales to zero
scale to zero: app engine standard, cloud function
upvoted 2 times

**vijbabu** 2 years, 5 months ago
Selected Answer: B
Answer is B
upvoted 1 times

**Dhiraj03** 2 years, 6 months ago
Selected Answer: A
Cloud Functions can scale to zero when not in use
upvoted 3 times

**Sskhan** 2 years, 11 months ago
Selected Answer: A
Answer is A, As Option D App engine flexible can have minimum 1 instance active.
upvoted 4 times

**Sskhan** 2 years, 11 months ago
Answer is A, As Option D App engine flexible can have minimum 1 instance active.
upvoted 1 times

Question #64 *Topic 1*

You are creating an App Engine application that uses Cloud Datastore as its persistence layer. You need to retrieve several root entities for which you have the identifiers. You want to minimize the overhead in operations performed by Cloud Datastore. What should you do?

A. Create the Key object for each Entity and run a batch get operation

B. Create the Key object for each Entity and run multiple get operations, one operation for each entity

C. Use the identifiers to create a query filter and run a batch query operation

D. Use the identifiers to create a query filter and run multiple query operations, one operation for each entity

**shashu07** Highly Voted 4 years ago
Correct Answer: A

Create the Key object for each Entity and run a batch get operation
https://cloud.google.com/datastore/docs/best-practices
Use batch operations for your reads, writes, and deletes instead of single operations. Batch operations are more efficient because they perfor multiple operations with the same overhead as a single operation.
Firestore in Datastore mode supports batch versions of the operations which allow it to operate on multiple objects in a single Datastore mod call.
Such batch calls are faster than making separate calls for each individual entity because they incur the overhead for only one service call. If multiple entity groups are involved, the work for all the groups is performed in parallel on the server side.

upvoted 48 times

**AzureDP900** 1 year, 8 months ago

works fine .. A is right

upvoted 2 times

**AWS56** Highly Voted 4 years, 5 months ago

Agree A

upvoted 7 times

**de1001c** Most Recent 3 weeks ago

Keep in mind that datastore is discontinued, Firestore being the recommended alternative.

upvoted 4 times

**don_v** 5 months, 1 week ago

According to
"A. Create the Key object for each Entity and run a batch get operation"
which is wrong as the key is already created for each entity whenever it's persisted.

I believe the correct answer is C, -- to use a bulk query (a.k.a. "batch" in their terms).
You need a query with a criteria anyway to find a resultset, and not just a fetch by "get" operation to load by surrogate keys.

upvoted 3 times

**vamgcp** 1 year, 4 months ago

By using the "lookup by key" API of Cloud Datastore, you can minimize the overhead in operations performed by Cloud Datastore and optimi the performance of your App Engine application.
from google.cloud import datastore

client = datastore.Client()
keys = [client.key('EntityKind', id) for id in entity_ids]
entities = client.get_multi(keys)

upvoted 3 times

**omermahgoub** 1 year, 6 months ago

A. Create the Key object for each Entity and run a batch get operation

To minimize the overhead in operations performed by Cloud Datastore, you should use the batch get operation to retrieve multiple entities in a single API call. To do this, you should create a Key object for each entity that you want to retrieve, then pass the Key objects to the batch get operation. This will allow you to retrieve multiple entities in a single API call, reducing the number of operations performed by Cloud Datastore and improving the efficiency of your application.

upvoted 6 times

**omermahgoub** 1 year, 6 months ago

Option B, running multiple get operations, one operation for each entity, would not be an efficient way to retrieve the entities because it wo require multiple API calls to Cloud Datastore, which would increase the overhead and decrease the efficiency of the application.

Option C, using the identifiers to create a query filter and running a batch query operation, would not be an efficient way to retrieve the ent because it would require performing a query operation, which is generally more expensive than a get operation.

Option D, using the identifiers to create a query filter and running multiple query operations, one operation for each entity, would not be an efficient way to retrieve the entities because it would require performing multiple query operations, which are generally more expensive tha get operations.

upvoted 5 times

**megumin** 1 year, 7 months ago

Selected Answer: A

A is ok

upvoted 1 times

**Mahmoud_E** 1 year, 8 months ago

**Selected Answer: A**

A is correct https://cloud.google.com/datastore/docs/best-practices#api_calls

upvoted 2 times

**RitwickKumar** 1 year, 10 months ago

**Selected Answer: A**

https://cloud.google.com/datastore/docs/concepts/entities#datastore-datastore-batch-lookup-python

upvoted 2 times

**haroldbenites** 2 years, 6 months ago

go for A.

upvoted 2 times

**vincy2202** 2 years, 6 months ago

A is the right answer

upvoted 2 times

**joe2211** 2 years, 6 months ago

**Selected Answer: A**

vote A

upvoted 2 times

**MaxNRG** 2 years, 7 months ago

A – create a key object for each entity, and run a batch get operations.
See Batch Operations section here: https://cloud.google.com/datastore/docs/concepts/entities
var keys = new Key[] { _keyFactory.CreateKey(1), _keyFactory.CreateKey(2) };
var tasks = _db.Lookup(keys[0], keys[1]);

1 and 2 are identifiers of the Key. Check Key / Identifier definition on the same link (top of that page)
Such batch calls are faster than making separate calls for each individual entity because they incur the overhead for only one service call.

upvoted 1 times

**victory108** 3 years, 1 month ago

A. Create the Key object for each Entity and run a batch get operation

upvoted 1 times

**un** 3 years, 1 month ago

A is correct

upvoted 1 times

**Ausias18** 3 years, 2 months ago

Answer is A

upvoted 1 times

**Ausias18** 3 years, 2 months ago

Answer is A

upvoted 1 times

Question #65 *Topic 1*

You need to upload files from your on-premises environment to Cloud Storage. You want the files to be encrypted on Cloud Storage using customer-supplied encryption keys. What should you do?