SureshbabuK Most Recent 6 months, 3 weeks ago

Selected Answer: B

Access to specific ports and protocol can be controlled only by firewall rule - Hence B is correct. D is not correct as service account is to authenticate and Authorized a specific machine to resource or service not ports and protocols

upvoted 4 times

megumin 7 months, 2 weeks ago

Selected Answer: B

B is ok

upvoted 1 times

AzureDP900 8 months ago

B is the best option.

upvoted 1 times

abirroy 9 months, 1 week ago

Selected Answer: B

Use firewall rules based on network tags attached to the compute instances

upvoted 2 times

alexandercamachop 9 months, 2 weeks ago

The secret is "paths and ports".
Which tell us Firewall as our only option.

upvoted 5 times

medi01 1 month, 4 weeks ago

And how does firewall restrict "paths" pretty please?

upvoted 3 times

cbarg 11 months, 2 weeks ago

Selected Answer: B

B Firewall rules to restrict traffic

upvoted 1 times

haroldbenites 1 year, 6 months ago

Go for B.

upvoted 2 times

vincy2202 1 year, 6 months ago

B is the right answer

upvoted 2 times

MaxNRG 1 year, 7 months ago

B – use firewall rules based on network tags attached to the compute instances
This answer avoids using IP, which are replaced by tags.

upvoted 3 times

MamthaSJ 1 year, 11 months ago

Answer is B

upvoted 4 times

areza 2 years ago

B is ok

upvoted 1 times

victory108 2 years, 1 month ago

B. Use firewall rules based on network tags attached to the compute instances

upvoted 2 times

Ausias18 2 years, 2 months ago

Answer is B

upvoted 1 times

lynx256 2 years, 2 months ago
B is ok
upvoted 1 times

Vika 2 years, 3 months ago
Agree B
upvoted 1 times

Question #73 *Topic 1*

You are using Cloud SQL as the database backend for a large CRM deployment. You want to scale as usage increases and ensure that you don't run out of storage, maintain 75% CPU usage cores, and keep replication lag below 60 seconds. What are the correct steps to meet your requirements?

A. 1. Enable automatic storage increase for the instance. 2. Create a Stackdriver alert when CPU usage exceeds 75%, and change the instance type to reduce CPU usage. 3. Create a Stackdriver alert for replication lag, and shard the database to reduce replication time.

B. 1. Enable automatic storage increase for the instance. 2. Change the instance type to a 32-core machine type to keep CPU usage below 75%. 3. Create a Stackdriver alert for replication lag, and deploy memcache to reduce load on the master.

C. 1. Create a Stackdriver alert when storage exceeds 75%, and increase the available storage on the instance to create more space. 2. Deploy memcached to reduce CPU load. 3. Change the instance type to a 32-core machine type to reduce replication lag.

D. 1. Create a Stackdriver alert when storage exceeds 75%, and increase the available storage on the instance to create more space. 2. Deploy memcached to reduce CPU load. 3. Create a Stackdriver alert for replication lag, and change the instance type to a 32-core machine type to reduce replication lag.

AWS56 Highly Voted 4 years, 11 months ago
Agree with A
upvoted 26 times

heretolearnazure 1 year, 3 months ago
Sharding database will reduce latency
upvoted 4 times

AzureDP900 2 years, 2 months ago
1. Enable automatic storage increase for the instance. 2. Create a Stackdriver alert when CPU usage exceeds 75%, and change the instan type to reduce CPU usage. 3. Create a Stackdriver alert for replication lag, and shard the database to reduce replication time.
upvoted 2 times

9xnine Highly Voted 2 years, 6 months ago
Has anyone who has taken the exam recently seen any lingering questions with the Stackdriver nomenclature or is it all cloud logging, cloud monitoring, etc.?
upvoted 16 times

nareshthumma Most Recent 1 month, 3 weeks ago
Answer is: A
upvoted 1 times

e5019c6 11 months, 3 weeks ago

Selected Answer: B

While I understand the doubts of selecting a 32 core machine from the start, answer A might be wrong...
According to this article:
https://cloud.google.com/sql/docs/mysql/instance-settings#impact
''For MySQL instances, changing either the machine type or the zone of the instance results in the instance going offline for several minutes.''
And I understand that instance type == machine type.
If we switch to a PostgreSQL or SQL Server instance, similar warnings appear:
PostgreSQL: ''Changing the number of CPUs or the memory size results in the instance going offline for less than 60 seconds. The total time the changes to take effect can take several minutes.''
SQL Server: ''Changing the number of CPUs or the memory size results in the instance going offline for less than 60 seconds.''

upvoted 2 times

e5019c6 11 months, 3 weeks ago

But I also understand that the question didn't mention anything about needing the DB to be online always, so maybe these offline times ar acceptable...

upvoted 1 times

AdityaGupta 1 year, 2 months ago

Selected Answer: A

You are using Cloud SQL as the database backend for a large CRM deployment. You want to scale as usage increases and ensure that you d run out of storage, maintain 75% CPU usage cores, and keep replication lag below 60 seconds. What are the correct steps to meet your requirements?

C & D is out of question as it is talking of 75% of storage, where in question it says 75% of CPU.
Option A says monitoring before before taking action and sharding will also help in reducing latency.
Option B specifies specific machine type, which is not correct and also memcache which is used to recude the round trip to fetch data, it will help in reducing latency.

I would prefer to go with Option A, as it is correct sequence to solve the problem.

upvoted 2 times

red_panda 1 year, 6 months ago

Selected Answer: A

For me is A.

upvoted 1 times

jfricker 1 year, 10 months ago

The correct answer is D.

1. Create a Stackdriver alert when storage exceeds 75%, and increase the available storage on the instance to create more space.
2. Deploy memcached to reduce CPU load.
3. Create a Stackdriver alert for replication lag, and change the instance type to a 32-core machine type to reduce replication lag.

This approach ensures that you are able to address the three requirements specified in the question:

- Monitoring storage usage and increasing storage when it exceeds 75% to avoid running out of storage.
- Reducing CPU load by deploying memcached, which can be used to cache frequently-used data, offloading some of the load from the database.
- Monitoring replication lag and increasing the number of cores to reduce lag.

upvoted 4 times

Charsoft 1 year, 11 months ago

It may be A for the simple fact that all the other answers throw in a tiny detail about 32 cores. This seems like a red herring (unnecessary deta that are meant to distract), so for that reason, A is the answer.

upvoted 5 times

megumin 2 years, 1 month ago

Selected Answer: A

A is ok

upvoted 1 times

AzureDP900 2 years, 2 months ago

A is right

upvoted 1 times

6721sora 2 years, 3 months ago

A is incorrect. because of the wording "Shard the database". How can you shard the database in Cloud SQL without causing major disruptior Sharding is not a core feature of RDBMS.

B should be correct. inspite of the mention of a fixed 32 core

upvoted 3 times

jay9114 2 years, 2 months ago

Ii
Kk
Kk
You can shard cloudsql. Review this article - https://cloud.google.com/community/tutorials/horizontally-scale-mysql-database-backend-w google-cloud-sql-and-proxysql#:~:text=Common%20approaches%20for%20horizontally%20scaling,with%20Cloud%20SQL%20and%20ProxySQL.

upvoted 4 times

fiercedog 2 years ago

The article only mentions sharding as a concept, and not a solution for cloudsql.

upvoted 1 times

Deb2293 1 year, 9 months ago

https://cloud.google.com/community/tutorials/horizontally-scale-mysql-database-backend-with-google-cloud-sql-and-proxysql#:~:text=SQL%20and%20ProxySQL.-,Sharding,logic%20or%20a%20query%20router.

You can shard MySQL.

Answer should be A.

upvoted 4 times

mj20201 2 years, 9 months ago

Selected Answer: A

vote for A

upvoted 2 times

haroldbenites 3 years ago

Go for A

upvoted 2 times

vincy2202 3 years, 1 month ago

A is the correct answer

upvoted 2 times

amxexam 3 years, 3 months ago

We can directly eliminate C and D we are doing some work that is already automated.

Still, I cannot make a point why not B is better than A?
I believe adding memcash will give an additional boost

Can someone help me point out why A is better than B?

upvoted 4 times

[Removed] 3 years, 2 months ago

Just to back up what amxexam said, here is the link on automatically increasing storage based on trend analysis:

https://cloud.google.com/sql/docs/mysql/instance-settings#storage-capacity-2ndgen

upvoted 3 times

HenkH 2 years, 11 months ago

That is correct - but doc only mentions auto storage increase for this specific product (cloud SQL).

upvoted 1 times

HenkH 2 years, 2 months ago

Should read MySQl

upvoted 2 times

cotam 3 years, 2 months ago

I suppose B is not a better option, since it indicates 'add 32core cpu', with no info of the current usage that seems like a over-kill.

upvoted 5 times

Ishu_awsguy 2 years, 4 months ago

I would say only because of the below line
"You want to scale as usage increases" Line 1
Creating a 32 core machine upfront where we do not know what was the source machine cores would not be ideal .
in that situation i would go with A

upvoted 4 times

MamthaSJ 3 years, 5 months ago

Answer is A

upvoted 4 times

areza 3 years, 6 months ago

A it is

upvoted 1 times

Question #74 *Topic 1*

You are tasked with building an online analytical processing (OLAP) marketing analytics and reporting tool. This requires a relational database that can operate on hundreds of terabytes of data. What is the Google-recommended tool for such applications?

A. Cloud Spanner, because it is globally distributed

B. Cloud SQL, because it is a fully managed relational database

C. Cloud Firestore, because it offers real-time synchronization across devices

D. BigQuery, because it is designed for large-scale processing of tabular data

AWS56 Highly Voted 4 years, 11 months ago
Agree D
upvoted 20 times

tartar 4 years, 4 months ago
D is ok
upvoted 5 times

Nastrand 3 years, 10 months ago
What about the relational part? BigQuery uses SQL but it's not relational... I'm not sure its D
upvoted 4 times

riflerrick 3 years, 6 months ago
BigQuery is relational!
upvoted 5 times

lovingsmart2000 3 years, 5 months ago
Pls do not confuse - Cloud SQL and BigQuery are RDBMS. Cloud Datastore, Bigtable are NoSQL.
Right answer is D - BQ
upvoted 14 times

kumarp6 4 years, 1 month ago
Yes it is D
upvoted 2 times

nitinz 3 years, 9 months ago
D, OLAP=BQ
upvoted 4 times

Sur_Nikki 1 year, 7 months ago
Well Said
upvoted 1 times

JasonL_GCP 3 years, 2 months ago
The question asks "This requires a relational database that can operate on hundreds of terabytes of data", but bq doesn't meet this condition?
upvoted 2 times

elaineshi 2 years, 6 months ago
BigQuery supports relational and query of join tables.
upvoted 2 times

gfhbox0083 Highly Voted 4 years, 6 months ago
D, for sure.
BigQuery for OLAP
Google Cloud Spanner for OLTP.
upvoted 15 times

gracjanborowiak Most Recent 5 months, 1 week ago

Selected Answer: D

sql is oltp
olap is data cube. lots of data which we try to process somehow. biq query is for that.

firestore is for mobile and web apps. not so fast nosql db.

upvoted 1 times

Anandmrk 10 months ago

Agreed

upvoted 1 times

AdityaGupta 1 year, 2 months ago

Selected Answer: D

4 reasons to choose BQ (Supports Petabytes of data)
- OLAP Data
- Relational DB (SQL)
- 100s of TB data
- Analystics and Reporting

upvoted 1 times

Ashish1995 1 year, 9 months ago

Selected Answer: D

D is obvious

upvoted 1 times

CGS22 1 year, 9 months ago

Selected Answer: D

D is the right one

upvoted 1 times

SudhirAhirkar 2 years ago

Cloud SQL/Spanner is OLTP DB but not OLAP. BQ is a well-known OLAP for analytics and also supports RBMS feature too... so I would got v
D

upvoted 1 times

AniketD 2 years ago

Selected Answer: D

D is correct. BigQuery is relational. Cloud SQL is not OLAP; moreover it can not store/process hundreds of TB of data. Max size is 64 TB only

upvoted 1 times

megumin 2 years, 1 month ago

Selected Answer: D

D is ok

upvoted 1 times

SerGCP 2 years, 1 month ago

Selected Answer: D

https://cloud.google.com/products/databases.

upvoted 1 times

zr79 2 years, 2 months ago

The words you need to focus "You are tasked with building an online analytical processing (OLAP) marketing analytics and reporting tool" whi
is BigQuery

upvoted 1 times

AzureDP900 2 years, 2 months ago

Big Query for large analytics , D is right

upvoted 1 times

Andre777 2 years, 3 months ago

Selected Answer: D

The keyword in this context is OLAP. CloudSQL is Relational SQL for OLTP. Capacity wise, BQ supports for PB+ while CloudSQL only have m capacity of up to ~10TB. Again the questions specifically mention "hundreds of TB of data". So D is the answer.

upvoted 2 times

deepdowndave 2 years, 3 months ago

Why is it not CloudSQL? It supports TB data storage and the question is about a relational database, not a data warehouse such as BigQuery

upvoted 1 times

Andre777 2 years, 3 months ago

The keyword in this context is OLAP. CloudSQL is Relational SQL for OLTP. Capacity wise, BQ supports for PB+ while CloudSQL only have max capacity of up to ~10TB. Again the questions specifically mention "hundreds of TB of data". So D is the answer.

upvoted 1 times

alexandercamachop 2 years, 3 months ago

The answer is Big Query, D
Secret: Analytical, Hundreds of TBTs. Relational.
All of this are strictly meet by Big Query, if it had not said Analytical but rather, other keywords like High Availability then Cloud Spanner.

upvoted 1 times

Thornadoo 2 years, 5 months ago

Selected Answer: D

Guys, this is easy:
OLTP - Cloud Spanner & Cloud SQL
OLAP - Big Query
NoSQL - Filestore and Big Table

So answer is D.

upvoted 1 times

Question #75 *Topic 1*

You have deployed an application to Google Kubernetes Engine (GKE), and are using the Cloud SQL proxy container to make the Cloud SQL database available to the services running on Kubernetes. You are notified that the application is reporting database connection issues. Your company policies require a post- mortem. What should you do?

A. Use gcloud sql instances restart.

B. Validate that the Service Account used by the Cloud SQL proxy container still has the Cloud Build Editor role.

C. In the GCP Console, navigate to Stackdriver Logging. Consult logs for (GKE) and Cloud SQL.

D. In the GCP Console, navigate to Cloud SQL. Restore the latest backup. Use kubectl to restart all pods.

jcmoranp Highly Voted 4 years, 1 month ago

post mortem always includes log analysis, answer is C

upvoted 65 times

Sur_Nikki 7 months, 2 weeks ago

Thanks for the info

upvoted 1 times

AzureDP900 1 year, 2 months ago

C is right for Root Cause Analysis.

upvoted 1 times

AWS56 3 years, 11 months ago
AGREE C
upvoted 3 times

MamthaSJ Highly Voted 2 years, 5 months ago
Answer is C
upvoted 5 times

pakilodi Most Recent 2 weeks, 4 days ago
Selected Answer: C
C -> post-mortem = log analysis
upvoted 2 times

AdityaGupta 2 months, 2 weeks ago
Selected Answer: C
You can jump on to the conlusion hence answer is not B. Consulting logs is always a good way to start investigation. and A and D is not a choice.
upvoted 1 times

ale_brd_111 1 year, 1 month ago
Stackdriver is deprecated, now you must navigate to Cloud Logging.
upvoted 2 times

megumin 1 year, 1 month ago
Selected Answer: C
C is ok
upvoted 1 times

Mahmoud_E 1 year, 1 month ago
Selected Answer: C
C is the right answer
upvoted 1 times

Jay_Krish 1 year, 3 months ago
Selected Answer: C
Logical answer is C. But is Stackdriver Logging enabled by default? Appreciate if someone could answer this?
upvoted 1 times

haroldbenites 2 years ago
Go for C
upvoted 1 times

pakilodi 2 years ago
Selected Answer: C
post mortem = logs
upvoted 1 times

vincy2202 2 years ago
C is the correct answer
upvoted 1 times

joe2211 2 years ago
Selected Answer: C
vote C
upvoted 1 times

MaxNRG 2 years, 1 month ago
C – in GCP Console navigate to Stackdriver Logging. Consult logs for Kubernetes Engine and Cloud SQL.
A/D – is an immediate attempt to fix an issue. No analysis.
B – is irrelevant at all. Cloud SQL proxy should not build anything in production.
upvoted 4 times

lovingsmart2000 2 years, 5 months ago

Answer is C. I request all here - not to blindly follow the answers published at coursera or udemy as most of them are copy-pasted answer an are not real. Examtopis provides the more accurate answers and also support with comments

upvoted 4 times

lovingsmart2000 2 years, 5 months ago

Answer is B. I request all here - not to blindly follow the answers published at coursera or udemy as most of them are copy-pasted answer an are not real. Examtopis provides the more accurate answers and also support with comments

upvoted 2 times

ashish_t 2 years, 2 months ago

Why Service Account needs Cloud Build Editor role for accessing Cloud SQL?
The role is misleading/wrong, so B is wrong.

upvoted 4 times

victory108 2 years, 7 months ago

C. In the GCP Console, navigate to Stackdriver Logging. Consult logs for Kubernetes Engine and Cloud SQL.

upvoted 2 times

un 2 years, 7 months ago

C is correct

upvoted 1 times

Question #76 *Topic 1*

Your company pushes batches of sensitive transaction data from its application server VMs to Cloud Pub/Sub for processing and storage. What is the Google- recommended way for your application to authenticate to the required Google Cloud services?

A. Ensure that VM service accounts are granted the appropriate Cloud Pub/Sub IAM roles.

B. Ensure that VM service accounts do not have access to Cloud Pub/Sub, and use VM access scopes to grant the appropriate Cloud Pub/Sub IAM roles.

C. Generate an OAuth2 access token for accessing Cloud Pub/Sub, encrypt it, and store it in Cloud Storage for access from each VM.

D. Create a gateway to Cloud Pub/Sub using a Cloud Function, and grant the Cloud Function service account the appropriate Cloud Pub/Sub IAM roles.

AWS56 Highly Voted 4 years, 5 months ago

Agree A

upvoted 26 times

nitinz 3 years, 3 months ago

A is correct

upvoted 2 times

kumarp6 3 years, 7 months ago

Yes A it is

upvoted 2 times

JustJack21 Highly Voted 2 years, 9 months ago

It's because of questions like these that I do not feel guilty about using question banks :D In what world would you accept value requirements this from your user? Wouldn't you ask "Do you want to just authenticate? or the data to be encrypted on its way to pub/sub?"
I'll ignore the first part of the question and assume all data is sensitive, and focus on "What is the Google- recommended way for your applica to authenticate to the required Google Cloud services?" -- The answer then is A.

Use encryption and defense-in-depth for the first part.

upvoted 12 times

bandegg 5 months, 2 weeks ago

> It's because of questions like these that I do not feel guilty about using question banks :D

Same. To me, it wasn't clear whether the servers were in google or not due to the question about accessing google cloud. It was asked as the VMs were outside of google

upvoted 3 times

AMEJack 1 year, 8 months ago

Service accounts use keys

upvoted 1 times

red_panda Most Recent 1 year ago

Selected Answer: A

A is correct for me. It's batch, so no cloud function

upvoted 2 times

omermahgoub 1 year, 6 months ago

A. Ensure that VM service accounts are granted the appropriate Cloud Pub/Sub IAM roles.

The Google-recommended way for your application to authenticate to Cloud Pub/Sub and other Google Cloud services when running on Compute Engine VMs is to use VM service accounts. VM service accounts are automatically created when you create a Compute Engine VM, and they are associated with the VM instance. To authenticate to Cloud Pub/Sub and other Google Cloud services, you should ensure that the VM service accounts are granted the appropriate IAM roles.

upvoted 7 times

omermahgoub 1 year, 6 months ago

Option B, ensuring that VM service accounts do not have access to Cloud Pub/Sub and using VM access scopes to grant the appropriate Cloud Pub/Sub IAM roles, would not be a suitable solution because VM service accounts are required for authentication to Google Cloud services.

Option C, generating an OAuth2 access token for accessing Cloud Pub/Sub, encrypting it, and storing it in Cloud Storage for access from each VM, would not be a suitable solution because it would require manual management of access tokens, which can be error-prone and insecure.

Option D, creating a gateway to Cloud Pub/Sub using a Cloud Function and granting the Cloud Function service account the appropriate Cloud Pub/Sub IAM roles, would not be a suitable solution because it would not allow the application to directly authenticate to Cloud Pub/Sub.

upvoted 3 times

Sur_Nikki 1 year, 1 month ago

Great way of explanation..By removing/elimination approach

upvoted 1 times

megumin 1 year, 7 months ago

Selected Answer: A

A is ok

upvoted 1 times

Mahmoud_E 1 year, 8 months ago

Selected Answer: A

A is the correct answer

upvoted 1 times

DrishaS4 1 year, 10 months ago

Selected Answer: A

https://cloud.google.com/iam/docs/understanding-service-accounts

upvoted 1 times

Pazzooo 2 years, 4 months ago

Selected Answer: A

The combination of Roles assigned to Service accounts granted to VMs is the way to go. :)

upvoted 2 times

elenamatay 2 years, 5 months ago

Service accounts are recommended for almost all cases in Pub/Sub (see https://cloud.google.com/pubsub/docs/authentication#service-accounts)

upvoted 3 times

haroldbenites 2 years, 6 months ago

Go for A.

upvoted 2 times

vincy2202 2 years, 6 months ago

A is the correct answer

upvoted 1 times

MaxNRG 2 years, 7 months ago

A – ensure that VM service accounts are granted the appropriate Cloud Pub/Sub IAM roles.
Check Migrating Data to GCP section of this page:
https://cloud.google.com/iam/docs/understanding-service-accounts
You will create a service account key and use it from an external process to call Cloud Platform APIs.

upvoted 3 times

Bakili 2 years, 8 months ago

A is very correct

upvoted 1 times

MamthaSJ 2 years, 11 months ago

Answer is A

upvoted 2 times

victory108 3 years, 1 month ago

A. Ensure that VM service accounts are granted the appropriate Cloud Pub/Sub IAM roles.

upvoted 3 times

AzureDP900 1 year, 8 months ago

Agreed with A

upvoted 1 times

un 3 years, 1 month ago

A is correct

upvoted 1 times

kartikjena31 3 years, 2 months ago

Ans. A

upvoted 1 times

Question #77 *Topic 1*

You want to establish a Compute Engine application in a single VPC across two regions. The application must communicate over VPN to an on-premises network.
How should you deploy the VPN?

A. Use VPC Network Peering between the VPC and the on-premises network.

B. Expose the VPC to the on-premises network using IAM and VPC Sharing.

C. Create a global Cloud VPN Gateway with VPN tunnels from each region to the on-premises peer gateway.

D. Deploy Cloud VPN Gateway in each region. Ensure that each region has at least one VPN tunnel to the on-premises peer gateway.

Googler2 Highly Voted 4 years, 8 months ago

It can't be -A - VPC Network Peering only allows private RFC 1918 connectivity across two Virtual Private Cloud (VPC) networks. In this exam is one VPC with on-premise network
https://cloud.google.com/vpc/docs/vpc-peering

It is not definitely - B - Can't be

It is not C - Because Cloud VPN gateways and tunnels are regional objects, not global

So, it the answer is D -
https://cloud.google.com/vpn/docs/how-to/creating-static-vpns

upvoted 45 times

amxexam 3 years, 3 months ago

Why not A?
https://cloud.google.com/vpc/docs/vpc-peering#benefits_of_exchanging_custom_routes
The second use case is exactly what is in the question.

Don't get the argument about RFC 1918.
Will go with A

upvoted 1 times

ochanz 3 years ago

https://cloud.google.com/vpc/docs/vpc-peering allows internal IP address connectivity across two VPC so A is not the answer as the o premise network need to use public IP. cmiiw

upvoted 4 times

AdityaGupta 1 year, 2 months ago

The question clearly asks us to use VPN.

upvoted 2 times

AzureDP900 2 years, 2 months ago

Agreed with D.

upvoted 1 times

TaherShaker Highly Voted 4 years, 1 month ago

Just Passed my exam and I answered (D) for this question

upvoted 20 times

M_Asep 3 years ago

sound promising dude

upvoted 3 times

Sur_Nikki 1 year, 7 months ago

IS the Exam Idea questions enough dude, for passing this exam?

upvoted 3 times

ccpmad Most Recent 5 months, 3 weeks ago

Selected Answer: D

Option C: Create a global VPN gateway and establish VPN tunnels from each region to the on-premises peer gateway. This suggests that a si
global VPN gateway manages the tunnels from both regions.
Option D: Deploy a VPN gateway in each region and ensure that each region has at least one VPN tunnel to the on-premises peer gateway. Th
indicates that each region has its own VPN gateway.

>Option D ensures that there is a VPN gateway in each region, providing greater redundancy. If a gateway in one region fails, the gateway in t
other region remains operational.

upvoted 1 times

santoshchauhan 8 months, 3 weeks ago

Selected Answer: C

Global Cloud VPN Gateway: This feature allows for the creation of a single VPN gateway that can serve multiple regions within the same VPC
network. By creating a global VPN gateway, you can efficiently manage VPN connections from all regions of your VPC to your on-premises
network.

Simplicity and Efficiency: Using a global gateway simplifies the configuration and management of VPN connections as opposed to maintainin
separate regional VPN gateways. It centralizes the VPN endpoint on the Google Cloud side, reducing the complexity of the network setup.

Reliable and Secure Communication: The global Cloud VPN Gateway allows for secure, encrypted tunnels between Google Cloud and the on
premises network, ensuring that the application's inter-regional and on-premises communications are secure.

upvoted 2 times

salvo007 11 months, 2 weeks ago

Selected Answer: D

C is wrong. A global vpn is a single region resource.
https://cloud.google.com/network-connectivity/docs/vpn/how-to/creating-ha-vpn?hl=it

gcloud compute vpn-gateways create GW_NAME \
--network=NETWORK \
--region=REGION \
--stack-type=IP_STACK

so D is the answer

upvoted 1 times

gcmrjbr 11 months, 3 weeks ago

It´s option C! So, while the VPN Gateway itself is a regional resource, its scope can be effectively global as it can serve resources across differ
regions within the same Virtual Private Cloud (VPC). This is why it's sometimes referred to as a 'global' service in the context of its functionalit
even though strictly speaking, it's a regional resource.

upvoted 2 times

AdityaGupta 1 year, 2 months ago

Selected Answer: D

Each Cloud VPN gateway is a regional resource that uses one or more regional external IP addresses. A Cloud VPN gateway can connect to a
peer VPN gateway.

upvoted 2 times

LaxmanTiwari 1 year, 7 months ago

It can't be -A - VPC Network Peering only allows private RFC 1918 connectivity across two Virtual Private Cloud (VPC) networks. In this exam is one VPC with on-premise network https://cloud.google.com/vpc/docs/vpc-peering It is not definitely - B - Can't be It is not C - Because Cl VPN gateways and tunnels are regional objects, not global So, it the answer is D - https://cloud.google.com/vpn/docs/how-to/creating-static-

upvoted 3 times

vvkds 1 year, 11 months ago

Selected Answer: D

D looks fine.

upvoted 1 times

oms_muc 2 years ago

Selected Answer: D

As HA isn't required, why do we need two VPN gateways?

upvoted 2 times

megumin 2 years, 1 month ago

Selected Answer: D

D is ok

upvoted 1 times

Mahmoud_E 2 years, 1 month ago

Selected Answer: D

D is the correct answer, in order to do A you will need VPN., or interconnect

upvoted 1 times

zr79 2 years, 2 months ago

there is two VPN:
1. classic VPN
2. HA VPN

upvoted 1 times

DrishaS4 2 years, 4 months ago

Selected Answer: D

Cloud VPN Gateway is a regional service, not global.

upvoted 4 times

elaineshi 2 years, 6 months ago

Why not C? services across regions can communicate to each other, VPN only connects to the closet region, and all the VPC shall be connec if firewall's set.

upvoted 2 times

haroldbenites 3 years ago

Go for D.
Cloud VPN Gateway is regional. NOt Global
gcloud compute vpn-gateways create GW_NAME \
--network=NETWORK \
--region=REGION

upvoted 2 times

vincy2202 3 years ago

D is the correct answer

upvoted 2 times

Question #78 *Topic 1*

Your applications will be writing their logs to BigQuery for analysis. Each application should have its own table. Any logs older than 45 days should be removed.
You want to optimize storage and follow Google-recommended practices. What should you do?

A. Configure the expiration time for your tables at 45 days

B. Make the tables time-partitioned, and configure the partition expiration at 45 days

C. Rely on BigQuery's default behavior to prune application logs older than 45 days

D. Create a script that uses the BigQuery command line tool (bq) to remove records older than 45 days

KouShikyou Highly Voted 4 years, 8 months ago
Could you please help clarify? I think B is correct.
It looks like table will be deleted with option A.
https://cloud.google.com/bigquery/docs/managing-tables#updating_a_tables_expiration_time
When you delete a table, any data in the table is also deleted. To automatically delete tables after a specified period of time, set the default tal expiration for the dataset or set the expiration time when you create the table.
upvoted 39 times

AzureDP900 1 year, 8 months ago
Agreed and going with B
upvoted 2 times

kumarp6 3 years, 7 months ago
it is B, if you use option A, on 46th day there is no table/content in table for application :)
upvoted 11 times

nitinz 3 years, 3 months ago
B partition table
upvoted 4 times

tartar 3 years, 10 months ago
B is ok
upvoted 8 times

aviv Highly Voted 4 years, 6 months ago
Agreed with B.
upvoted 10 times

OSAMA911 Most Recent 3 months, 4 weeks ago
Selected Answer: B
I think B is correct.
upvoted 1 times

AdityaGupta 8 months, 2 weeks ago
Selected Answer: B
https://cloud.google.com/bigquery/docs/managing-partitioned-tables#partition-expiration

B is the correct answer.
upvoted 3 times

SSPPJi 11 months, 2 weeks ago
https://cloud.google.com/bigquery/docs/managing-partitioned-tables#partition-expiration
upvoted 4 times

FaizAhmed 12 months ago
Selected Answer: B
B is correct
upvoted 1 times

Sur_Nikki 1 year, 1 month ago

B seems correct as this will partitioning will create a filter criteria on the basis of which specified actions on logs will be taken

upvoted 1 times

examch 1 year, 5 months ago

Selected Answer: B

B is the correct answer,

If your tables are partitioned by date, the dataset's default table expiration applies to the individual partitions. You can also control partition expiration using the time_partitioning_expiration flag in the bq command-line tool or the expirationMs configuration setting in the API. When a partition expires, data in the partition is deleted but the partitioned table is not dropped even if the table is empty.

https://cloud.google.com/bigquery/docs/best-practices-storage

upvoted 6 times

megumin 1 year, 7 months ago

Selected Answer: B

B is ok

upvoted 1 times

MarcoEscanor 1 year, 7 months ago

Selected Answer: B

B - You can control partition expiration using the time_partitioning_expiration flag in the bq command-line
https://cloud.google.com/bigquery/docs/best-practices-storage

upvoted 2 times

AhmedH7793 1 year, 9 months ago

Selected Answer: B

B is okay

upvoted 1 times

DrishaS4 1 year, 10 months ago

Selected Answer: B

Using Table-Partitions.

upvoted 1 times

DrishaS4 1 year, 10 months ago

Using Table-Partitions.

upvoted 1 times

[Removed] 2 years, 4 months ago

Selected Answer: B

I got similar question on my exam.

upvoted 1 times

haroldbenites 2 years, 6 months ago

Go for B.
https://cloud.google.com/bigquery/docs/creating-partitioned-tables#sql
CREATE TABLE
mydataset.newtable (transaction_id INT64, transaction_date DATE)
PARTITION BY
transaction_date
OPTIONS(
partition_expiration_days=3,
require_partition_filter=true
)

upvoted 2 times

vincy2202 2 years, 6 months ago

Selected Answer: B

B is the correct answer

upvoted 2 times

MaxNRG 2 years, 7 months ago

B – Make the tables time-partitioned and configure the partition expiration at 45 days.
A – if you use table expiration time, then it will remove the whole table after 45 days.
D – requires extra work and is not automatic.

upvoted 2 times

Question #79 *Topic 1*

You want your Google Kubernetes Engine cluster to automatically add or remove nodes based on CPU load.
What should you do?

A. Configure a HorizontalPodAutoscaler with a target CPU usage. Enable the Cluster Autoscaler from the GCP Console.

B. Configure a HorizontalPodAutoscaler with a target CPU usage. Enable autoscaling on the managed instance group for the cluster using the gcloud command.

C. Create a deployment and set the maxUnavailable and maxSurge properties. Enable the Cluster Autoscaler using the gcloud command.

D. Create a deployment and set the maxUnavailable and maxSurge properties. Enable autoscaling on the cluster managed instance group from the GCP Console.

Unfaithful Highly Voted 2 years, 5 months ago

Answer: A
Support:
How does Horizontal Pod Autoscaler work with Cluster Autoscaler?

Horizontal Pod Autoscaler changes the deployment's or replicaset's number of replicas based on the current CPU load. If the load increases, HPA will create new replicas, for which there may or may not be enough space in the cluster. If there are not enough resources, CA will try to bring up some nodes, so that the HPA-created pods have a place to run. If the load decreases, HPA will stop some of the replicas. As a result some nodes may become underutilized or completely empty, and then CA will terminate such unneeded nodes.

upvoted 62 times

heretolearnazure 3 months, 3 weeks ago

very well explained

upvoted 1 times

AzureDP900 1 year, 2 months ago

Nice and detailed explanation. I agree with A.

upvoted 1 times

LaxmanTiwari 7 months, 1 week ago

Nice and detailed explanation. I agree with A.

upvoted 1 times

Rajasa 2 years ago

Good Explaination

upvoted 3 times

natpilot Highly Voted 3 years, 10 months ago

i'm for A, but the question in ambiguous, because requires the autoscale of nodes (not pod) when the cpu overload, but in answer use k8s po autoscaler based on cpu load (cpu load for pod, not nodes). strange

upvoted 25 times

p4 3 years, 1 month ago

Agreed, the question is not about pods, but answers are also talking about pods (not only)
A is correct because B is wrong according to
https://cloud.google.com/kubernetes-engine/docs/concepts/cluster-autoscaler

"Caution: Do not enable Compute Engine autoscaling for managed instance groups for your cluster nodes. GKE's cluster autoscaler is separate from Compute Engine autoscaling"

upvoted 20 times

skywalker 3 years, 7 months ago

Confuse with the question like you mentioned. Autoscale is via nodes not pod.. and can only be configure using gcloud command.

upvoted 6 times

LaxmanTiwari Most Recent 7 months, 1 week ago

Nice and detailed explanation. I agree with A.

upvoted 1 times

Sur_Nikki 7 months, 2 weeks ago

A seems correct. y to create managed instance groups unnecessarily?

upvoted 1 times

Deb2293 9 months, 3 weeks ago

The answer is A.
More nodes mean it's horizontal scaling (increase VMs means vertical scaling of infrastructure). Cluster AutoScalar is used for increasing num of nodes.

upvoted 1 times

examch 11 months, 3 weeks ago

Selected Answer: A

A is the Correct answer, Horizontal Pod Autoscaler and Cluster Autoscaler can be used together to provision new pods and new nodes as pe CPU utilization.

https://www.youtube.com/watch?v=VNAWA6NkoBs

upvoted 2 times

megumin 1 year, 1 month ago

Selected Answer: A

ok for A

upvoted 1 times

Rajeev26 1 year, 2 months ago

Selected Answer: A

MIG not for GKE as option B and C, D are not relevant to question

upvoted 1 times

abirroy 1 year, 3 months ago

Selected Answer: A

Configure a HorizontalPodAutoscaler with a target CPU usage. Enable the Cluster Autoscaler from the GCP Console.

upvoted 1 times

gee1979 1 year, 3 months ago

Selected Answer: A

A...

The HPA and CA complement each other for truly efficient scaling. If the load increases, HPA will create new replicas. If there isn't enough spa for these replicas, CA will provision some nodes, so that the HPA-created pods have a place to run.

The Horizontal Pod Autoscaler changes the shape of your Kubernetes workload by automatically increasing or decreasing the number of Pod response to the workload's CPU or memory consumption, or in response to custom metrics reported from within Kubernetes or external metr from sources outside of your cluster.

upvoted 1 times

6721sora 1 year, 3 months ago

A is wrong.
Pod scaling only spins up additional pods. Not nodes.
Cluster Autoscaler does adding of nodes automatically.
I am surprised that so many people think that A is the correct answer.

Correct answer per me is C

upvoted 2 times

DrishaS4 1 year, 4 months ago

Selected Answer: A

Horizontal Pod Autoscaler changes the deployment's or replicaset's number of replicas based on the current CPU load. If the load increases, HPA will create new replicas, for which there may or may not be enough space in the cluster. If there are not enough resources, CA will try to bring up some nodes, so that the HPA-created pods have a place to run. If the load decreases, HPA will stop some of the replicas. As a result some nodes may become underutilized or completely empty, and then CA will terminate such unneeded nodes.

upvoted 2 times

[Removed] 1 year, 10 months ago

I got one question on my exam which showed autoscaling configuration and was asked to select correct configuration.

upvoted 1 times

OrangeTiger 1 year, 11 months ago

I agree A is correct.
I found quicklab.
Understanding and Combining GKE Autoscaling Strategies.

upvoted 1 times

ehgm 1 year, 11 months ago

Selected Answer: A

B and D: You must never change the GKE managed instance group.
C and D: maxUnavailable and maxSurge are used for rolling update
A. It is the correct.

upvoted 5 times

haroldbenites 2 years ago

Go for A

upvoted 1 times

Question #80 *Topic 1*

You need to develop procedures to verify resilience of disaster recovery for remote recovery using GCP. Your production environment is hosted on-premises. You need to establish a secure, redundant connection between your on-premises network and the GCP network.
What should you do?