examch Most Recent 5 months, 3 weeks ago

Selected Answer: B

B is the correct Answer,

Google Cloud HTTP(S) load balancers and Traffic Director use a Google Cloud configuration resource called a URL map to route HTTP(S) requests to backend services or backend buckets.

For example, with an external HTTP(S) load balancer, you can use a single URL map to route requests to different destinations based on the r configured in the URL map:

Requests for https://example.com/video go to one backend service.
Requests for https://example.com/audio go to a different backend service.
Requests for https://example.com/images go to a Cloud Storage backend bucket.
Requests for any other host and path combination go to a default backend service.
URL maps are used with the following Google Cloud products:

External HTTP(S) Load Balancing (global, regional, and classic modes)
Internal HTTP(S) Load Balancing
Traffic Director

https://cloud.google.com/load-balancing/docs/url-map-concepts

upvoted 6 times

omermahgoub 6 months ago

B. Create an HTTPS load balancer with URL Maps.

An HTTPS load balancer is a type of load balancer that can distribute incoming HTTPS traffic to one or more back-end services, such as Compute Engine instances or Google Kubernetes Engine clusters. It can also provide SSL/TLS termination, enabling you to use your own SSL/TLS certificates and keys.

You can use URL Maps to configure the HTTPS load balancer to route traffic based on the URL path being requested. This allows you to set u different URL paths to be served by different back-end services, providing a high level of flexibility in your load balancing configuration.

upvoted 1 times

omermahgoub 6 months ago

Option A, creating a cross-region load balancer with URL Maps, is also a valid solution, but it is not specifically designed for end-to-end in transit encryption.

Option C, creating appropriate instance groups and instances and configuring SSL proxy load balancing, is not a complete solution for glo load balancing. SSL proxy load balancing is a feature that enables you to terminate SSL/TLS connections at the load balancer and establis new SSL/TLS connection between the load balancer and the back-end service. It is not a global load balancing solution in and of itself.

Option D, creating a global forwarding rule and configuring SSL proxy load balancing, is not a complete solution for global load balancing based on the URL path being requested. A global forwarding rule is a type of load balancing configuration that directs traffic to a specific back-end service based on the IP address and port of the incoming request. It does not allow for routing based on the URL path.

Regenerate

upvoted 3 times

TonytheTiger 6 months, 2 weeks ago

Answer B: URL maps used with global external HTTP(S) load balancers and regional external HTTP(S) load balancer support several advance traffic management features such as header-based traffic steering, weight-based traffic splitting, and request mirroring.

https://cloud.google.com/load-balancing/docs/https#url-maps

upvoted 2 times

megumin 7 months, 1 week ago

Selected Answer: B

B is ok

upvoted 1 times

Sbgani 9 months, 2 weeks ago

Selected Answer: B

UrlMaps are used to route requests to a backend service based on rules that you define for the host and path of an incoming URL.

upvoted 5 times

Sbgani 9 months, 2 weeks ago

https://registry.terraform.io/providers/hashicorp/google/latest/docs/resources/compute_url_map ANS B

upvoted 2 times

DrishaS4 10 months, 2 weeks ago

Selected Answer: B

https://cloud.google.com/load-balancing/docs/https/url-map

upvoted 4 times

AzureDP900 8 months ago

thank you for pointing the link

upvoted 1 times

haroldbenites 1 year, 6 months ago

Go for B

upvoted 2 times

vincy2202 1 year, 6 months ago

Selected Answer: B

B is correct answer

upvoted 2 times

nqthien041292 1 year, 6 months ago

Selected Answer: B

Vote B

upvoted 1 times

un 2 years, 1 month ago

B is correct

upvoted 1 times

ccmcwolf 2 years, 2 months ago

there are interl https load balancers they are regional https://cloud.google.com/load-balancing/docs/l7-internal

upvoted 1 times

Ausias18 2 years, 2 months ago

Answer is B

upvoted 2 times

bnlcnd 2 years, 4 months ago

confused with A vs B. A has the word "cross region" but finally find out HTTP/S Load Balancing is naturally global.
- B

upvoted 5 times

doumx 2 years, 6 months ago

B easy

upvoted 2 times

Question #95 *Topic 1*

You have an application that makes HTTP requests to Cloud Storage. Occasionally the requests fail with HTTP status codes of 5xx and 429.
How should you handle these types of errors?

A. Use gRPC instead of HTTP for better performance.

B. Implement retry logic using a truncated exponential backoff strategy.

C. Make sure the Cloud Storage bucket is multi-regional for geo-redundancy.

D. Monitor https://status.cloud.google.com/feed.atom and only make requests if Cloud Storage is not reporting an incident.

bigob4ek Highly Voted 3 years, 7 months ago
Answer is B
You should use exponential backoff to retry your requests when receiving errors with 5xx or 429 response codes from Cloud Storage.
https://cloud.google.com/storage/docs/request-rate
upvoted 41 times

nitinz 2 years, 3 months ago
It is B
upvoted 1 times

AzureDP900 8 months ago
I agree with you, B should be right
upvoted 1 times

Sbgani Highly Voted 9 months, 2 weeks ago
HTTP 408, 429, and 5xx response codes.

Exponential backoff algorithm
For requests that meet both the response and idempotency criteria, you should generally use truncated exponential backoff.

Truncated exponential backoff is a standard error handling strategy for network applications in which a client periodically retries a failed reque with increasing delays between requests.

An exponential backoff algorithm retries requests exponentially, increasing the waiting time between retries up to a maximum backoff time. Se the following workflow example to learn how exponential backoff works:

You make a request to Cloud Storage.

If the request fails, wait 1 + random_number_milliseconds seconds and retry the request.

If the request fails, wait 2 + random_number_milliseconds seconds and retry the request.

If the request fails, wait 4 + random_number_milliseconds seconds and retry the request.

And so on, up to a maximum_backoff time.

Continue waiting and retrying up to a maximum amount of time (deadline), but do not increase the maximum_backoff wait period between ret
upvoted 13 times

omermahgoub Most Recent 6 months ago

. Implement retry logic using a truncated exponential backoff strategy.

HTTP status codes of 5xx and 429 typically indicate that there is a temporary issue with the service or that the rate of requests is too high. To handle these types of errors, it is generally recommended to implement retry logic in your application using a truncated exponential backoff strategy.

Truncated exponential backoff involves retrying the request after an initial delay, and then increasing the delay exponentially for each subsequ retry up to a maximum delay. This approach helps to reduce the number of failed requests and can improve the reliability of your application.

upvoted 3 times

omermahgoub 6 months ago

Option A, using gRPC instead of HTTP for better performance, is not directly related to handling HTTP status codes of 5xx and 429. gRPC a high-performance RPC framework that can be used in place of HTTP, but it is not a solution for handling errors.

Option C, making sure the Cloud Storage bucket is multi-regional for geo-redundancy, may help improve the reliability of the service, but it not a solution for handling errors.

Option D, monitoring https://status.cloud.google.com/feed.atom and only making requests if Cloud Storage is not reporting an incident, is a practical solution for handling errors. This approach would require constantly monitoring the status page and could result in significant delays in processing requests. Instead, it is generally recommended to implement retry logic in your application to handle errors.

upvoted 2 times

megumin 7 months, 1 week ago

Selected Answer: B

B is ok

upvoted 1 times

Sbgani 9 months, 2 weeks ago

Selected Answer: B

https://cloud.google.com/storage/docs/retry-strategy

upvoted 2 times

DrishaS4 10 months, 2 weeks ago

Selected Answer: B

2xx – successful requests;
4xx, 5xx – failed requests;
3xx – requests that require redirect.
https://cloud.google.com/storage/docs/json_api/v1/status-codes

upvoted 2 times

haroldbenites 1 year, 6 months ago

Go for B

upvoted 2 times

vincy2202 1 year, 6 months ago

B is the correct answer

upvoted 2 times

nqthien041292 1 year, 6 months ago

Selected Answer: B

Vote B

upvoted 2 times

joe2211 1 year, 6 months ago

Selected Answer: B

vote B

upvoted 1 times

MaxNRG 1 year, 7 months ago

B – Implement retry logic using a truncated exponential backoff strategy.
Per HTTP status and error codes for JSON the status codes are:
2xx – successful requests;
4xx, 5xx – failed requests;
3xx – requests that require redirect.
https://cloud.google.com/storage/docs/json_api/v1/status-codes
429 – Too many requests: your app tries to use more that its limit, additional requests will fail. Decrease your client's requests and/or use truncated exponential backoff (used for all requests with 5xx and 429 errors).
https://cloud.google.com/storage/docs/retry-strategy

upvoted 2 times

victory108 2 years, 1 month ago

B. Use Deployment Manager to automate service provisioning. Use Stackdriver to monitor and debug your tests.

upvoted 2 times

victory108 2 years, 1 month ago

This B. Implement retry logic using a truncated exponential backoff strategy.

upvoted 1 times

un 2 years, 1 month ago

Answer is B.
Link provided by bigob4ek has details

upvoted 1 times

Ausias18 2 years, 2 months ago

Answer is B

upvoted 1 times

CloudGenious 2 years, 4 months ago

As per google, if you run into any issue as increase latency or erroe rate ,pause your ramp up this give cloudstorage more time to scale your bucket . Best is backoff when 5xx ,429,408 response code

upvoted 2 times

bnlcnd 2 years, 4 months ago

https://cloud.google.com/storage/docs/exponential-backoff
- B

upvoted 2 times

awadheshk 2 years, 8 months ago

B is correct

upvoted 1 times

Question #96 *Topic 1*

You need to develop procedures to test a disaster plan for a mission-critical application. You want to use Google-recommended practices and native capabilities within GCP.
What should you do?

What should you do?

A. Use Deployment Manager to automate service provisioning. Use Activity Logs to monitor and debug your tests.

B. Use Deployment Manager to automate service provisioning. Use Stackdriver to monitor and debug your tests.

C. Use gcloud scripts to automate service provisioning. Use Activity Logs to monitor and debug your tests.

D. Use gcloud scripts to automate service provisioning. Use Stackdriver to monitor and debug your tests.

crypt0 Highly Voted 4 years, 1 month ago
I think answer B is correct:
https://cloud.google.com/solutions/dr-scenarios-planning-guide
upvoted 53 times

nitinz 2 years, 9 months ago
It is B, Google Best practice ---> never use scripts. They do not trust anyone else's code it seems.
upvoted 12 times

fraloca 2 years, 11 months ago
https://cloud.google.com/solutions/dr-scenarios-planning-guide#test_your_plan_regularly
upvoted 2 times

kumarp6 3 years, 1 month ago
B is correct
upvoted 2 times

tartar 3 years, 4 months ago
B is ok
upvoted 11 times

passnow Highly Voted 4 years ago
Boom, everyone studied and did their labs, stackdriver is google's recommended tool for monitoring and debbuging. I agree with u all that B i
the correct answer
upvoted 23 times

duzapo Most Recent 3 months, 1 week ago
Selected Answer: D
D its correct cause are 3 multiregions-availables and one bucket only can deploy in one multi region
https://cloud.google.com/storage/docs/locations?hl=es-419#location-mr
upvoted 1 times

jalberto 4 months ago
Selected Answer: B
I think B is the correct answer
upvoted 1 times

faridomu 6 months, 1 week ago
Why not A?
upvoted 2 times

Jlharidon 1 year ago
Selected Answer: B
Deploy managment + Stackdriver trained ig GCSB
upvoted 1 times

megumin 1 year, 1 month ago
Selected Answer: B
B is ok
upvoted 1 times

AzureDP900 1 year, 2 months ago
B is right
upvoted 1 times

tycho 1 year, 4 months ago
in practice, D could work as well..
upvoted 1 times

adacek1 10 months, 2 weeks ago
yeah, but only native solutions should be taken into consideration (as stated in requirements), so scripts are basically ruled out
upvoted 1 times

gaojun 1 year, 8 months ago
Answer B is correct
upvoted 1 times

Skr6266 1 year, 9 months ago
Selected Answer: B
Deployment Manager + Cloud Monitoring and Logging solution.
upvoted 1 times

haroldbenites 2 years ago
Go for B
upvoted 1 times

vincy2202 2 years ago
Selected Answer: B
B is the correct answer
upvoted 1 times

nqthien041292 2 years ago
Selected Answer: B
Vote B
upvoted 1 times

ganeshrev 2 years, 1 month ago
Selected Answer: B
Google recommended Practice
upvoted 1 times

victory108 2 years, 7 months ago
B. Use Deployment Manager to automate service provisioning. Use Stackdriver to monitor and debug your tests.
upvoted 4 times

un 2 years, 7 months ago
B is correct
upvoted 1 times

Question #97 *Topic 1*

Your company creates rendering software which users can download from the company website. Your company has customers all over the world. You want to minimize latency for all your customers. You want to follow Google-recommended practices.
How should you store the files?

A. Save the files in a Multi-Regional Cloud Storage bucket.

B. Save the files in a Regional Cloud Storage bucket, one bucket per zone of the region.

C. Save the files in multiple Regional Cloud Storage buckets, one bucket per zone per region.

D. Save the files in multiple Multi-Regional Cloud Storage buckets, one bucket per multi-region.

JoeShmoe Highly Voted 5 years, 1 month ago
Its D, create multi region buckets in Americas, Europe and Asia
upvoted 65 times

AmitAr 2 years, 7 months ago
What is point of Multi-Regional bucket, if this need to saved multiple times. I believe option (D) is for creating confusion only. It should be (A
upvoted 16 times

JaimeMS 6 months, 2 weeks ago
Let's try option A: you select a single multi-region bucket (e.g. Americas). Are you improving the latency of your clients in Asia? You do

Thus, Option A is not complete.
upvoted 2 times

AmitAr 2 years, 7 months ago
Read the question again.. I think (d) is correct.. eg. 1 bucket in US-multi-region, 2nd in AS-multi-region, 3rd in EU-multi-region
upvoted 10 times

giovanicascaes 1 year, 9 months ago
Yes, D seems correct. There are 3 multi-regions: ASIA, EU and US. In order to be global, there must be multi-region buckets in this 3 locations.

Reference: https://cloud.google.com/storage/docs/locations#location-mr
upvoted 6 times

turbo8p 2 years, 1 month ago
Check the current create bucket UI. You cannot select Asia multi-region and US multi-region at the same go. So to support global customer, you need to create multiple Multi-region buckets.
upvoted 13 times

Urban_Life 3 years ago
This can't be D. It should be A.
upvoted 6 times

kilo10x 1 year, 5 months ago
wrong its A
upvoted 1 times

MyPractice 4 years, 11 months ago
why " multiple Multi-Regional"? - A should be the right ans & addressing the global users - "More importantly, is that multiregional heavily leverages Edge caching and CDNs to provide the content to the end user"
https://medium.com/google-cloud/google-cloud-storage-what-bucket-class-for-the-best-performance-5c847ac8f9f2
upvoted 14 times

xavi1 3 years, 4 months ago
because a multi-regional includes all the locations of ONE region, not the others.
upvoted 13 times

MeasService Highly Voted 5 years, 2 months ago
I would go with A (https://cloud.google.com/storage/docs/locations)
upvoted 32 times

Amrx Most Recent 1 month, 3 weeks ago

Selected Answer: D

D is correct. Multi-region buckets are still specific to their own regional area, Americas, Europe and Asia. It's not A, doesn't cover the whole world.

upvoted 1 times

hehe_24 1 month, 4 weeks ago

I go for A. I never came across"multiple multi-region"

upvoted 1 times

JaimeMS 6 months, 2 weeks ago

Selected Answer: D

Its D, create multi region buckets in Americas, Europe and Asia

upvoted 2 times

tlopsm 6 months, 3 weeks ago

I think it is D.
Keyword: Your company has customers all over the world.
Lists Multi Regional cloud
Multi-Region Name Multi-Region Description
ASIA Data centers in Asia, excluding Hong Kong and Indonesia
EU Data centers within member states of the European Union*
US Data centers in the United States

Ans A. suggests "A" multi-regional Cloud. that means one of the above multi-regional cloud
Ans D: suggests "multiple" Multi-Regional so 2 or (preferably) all of the multi-regional cloud with one bucket per multi-region (less task)

upvoted 1 times

dija123 8 months, 2 weeks ago

Selected Answer: A

Agree with A

upvoted 1 times

5091a99 9 months, 1 week ago

Answer is A.
As for D: This would lead to data duplication and increased storage costs, as well as potential data consistency issues across different multi-regional buckets.

upvoted 1 times

Amrita2012 10 months ago

Selected Answer: A

Behind Multi-Regional Cloud Storage bucket, CDN is used hence good option to use for software download service.

upvoted 2 times

YasserM 10 months ago

Selected Answer: A

Buckets are not with zones, so B,C,D should be wrong. I go with A

upvoted 2 times

xaqanik 10 months, 2 weeks ago

I vote for A. Creating multiple multiregional Bucket is not seems practical and efficient. Besides that you can use CDN for minimizing latency a

upvoted 1 times

SSS987 11 months, 1 week ago

One doubt - if we are going with option D, how can we handle this in the application logic - pointing to different buckets depending on region'
Pls suggest.

upvoted 1 times

e5019c6 11 months, 3 weeks ago

Selected Answer: C

I'm going with C. Totally agree with 'theBestStudent' comment.
Multi-regional offers more availability for files, but worse latency, which was the requirement.
This link says it all:
https://cloud.google.com/storage/docs/locations#considerations
If you look at the table you can see the performance of the regional bucket (200Gbps) is much higher than the Multi-regional (50Gbps).
Also, Multi-regional buckets are 'only' available in US, Europe and East Asia. You would be letting behind places like South America, Africa, Canada, India, Indonesia, Middle East and Australia.

upvoted 1 times

theBestStudent 1 year ago

Selected Answer: C

It is C. Certainly is C, the other ones make no sense.

Multi regional is mostly for HA, performance is lower than regional. Regional gives better latency, so what you need to do is to have multiple regional buckets, in different regions of course https://cloud.google.com/storage/docs/locations

upvoted 4 times

yilexar 1 year, 2 months ago

A is correct. Regional and Dual-Regional buckets are optimized for latency. 200Gps vs. 50Gps of multi-region bucket.
https://cloud.google.com/storage/docs/locations

upvoted 2 times

Arun_m_123 1 year, 2 months ago

Selected Answer: A

A is the right answer- Multi-region buckets should be used for high availability / content delivery.

D is the wrong answer - There is nothing called "multiple multi region". This coinage itself is wrong

upvoted 2 times

steghe 1 year, 2 months ago

Selected Answer: D

D is correct. https://cloud.google.com/storage/docs/locations#location-mr

upvoted 1 times

Question #98 *Topic 1*

Your company acquired a healthcare startup and must retain its customers' medical information for up to 4 more years, depending on when it was created. Your corporate policy is to securely retain this data, and then delete it as soon as regulations allow.
Which approach should you take?

A. Store the data in Google Drive and manually delete records as they expire.

B. Anonymize the data using the Cloud Data Loss Prevention API and store it indefinitely.

C. Store the data in Cloud Storage and use lifecycle management to delete files when they expire.

D. Store the data in Cloud Storage and run a nightly batch script that deletes all expired data.

AWS56 Highly Voted 4 years, 11 months ago

Agree C

upvoted 23 times

desertlotus1211 Most Recent 3 weeks, 2 days ago

Selected Answer: B

Why not B? It's patients' health records...

upvoted 1 times

i_maddog_i 1 year, 9 months ago

Selected Answer: C

It's C

upvoted 1 times

megumin 2 years, 1 month ago

Selected Answer: C

C is ok

upvoted 1 times

AzureDP900 2 years, 2 months ago

I agree with C

upvoted 1 times

ACE_ASPIRE 2 years, 4 months ago

I got this question in exam.

upvoted 4 times

DrishaS4 2 years, 4 months ago

Selected Answer: C

go for C

upvoted 1 times

Dhiraj03 2 years, 6 months ago

Options C undoubtedly

upvoted 1 times

gaojun 2 years, 8 months ago

Go for C

upvoted 1 times

[Removed] 2 years, 10 months ago

I got similar question on my exam which involved life cycle management and bucket lock.

upvoted 3 times

Rajasa 2 years, 11 months ago

Selected Answer: C

Go for C

upvoted 1 times

haroldbenites 3 years ago

Go for C

upvoted 1 times

vincy2202 3 years ago

Selected Answer: C

C is the correct answer

upvoted 1 times

gabrielzeven 3 years ago

D sounds like i would do it, but C sound like a lab or exam

upvoted 1 times

nqthien041292 3 years ago

Selected Answer: C

Vote C

upvoted 1 times

victory108 3 years, 7 months ago

C. Store the data in Cloud Storage and use lifecycle management to delete files when they expire.

upvoted 4 times

un 3 years, 7 months ago
C is correct
upvoted 1 times

Question #99 *Topic 1*

You are deploying a PHP App Engine Standard service with Cloud SQL as the backend. You want to minimize the number of queries to the database.
What should you do?

A. Set the memcache service level to dedicated. Create a key from the hash of the query, and return database values from memcache before issuing a query to Cloud SQL.

B. Set the memcache service level to dedicated. Create a cron task that runs every minute to populate the cache with keys containing query results.

C. Set the memcache service level to shared. Create a cron task that runs every minute to save all expected queries to a key called ɹ€cached_queriesɹ€.

D. Set the memcache service level to shared. Create a key called ɹ€cached_queriesɹ€, and return database values from the key before using a query to Cloud SQL.

hiteshrup Highly Voted 3 years ago
A dedicated memset is always better than shared until cost-effectiveness specify in the exam as objective. So Option C and D are ruled out.

From A and B, Option B is sending and updating query every minutes which is over killing. So reasonable option left with A which balance performance and cost.

My answer will be A
upvoted 28 times

ArtistS 1 month ago
Good job bro
upvoted 1 times

Eroc Highly Voted 4 years, 1 month ago
https://cloud.google.com/appengine/docs/standard/php/memcache/using
upvoted 23 times

nitinz 2 years, 9 months ago
A is correct
upvoted 6 times

dlzhang 2 years, 6 months ago
https://cloud.google.com/memorystore/docs/redis/redis-overview
upvoted 2 times

tartar 3 years, 4 months ago
A is ok
upvoted 11 times

Sur_Nikki Most Recent 7 months, 2 weeks ago
Best is A
upvoted 2 times

megumin 1 year, 1 month ago
Selected Answer: A
A is ok
upvoted 1 times

AzureDP900 1 year, 2 months ago
A is fine.. dedicated mem cache
upvoted 1 times

ACE_ASPIRE 1 year, 4 months ago
I got this question in exam.
upvoted 5 times

DrishaS4 1 year, 4 months ago
Selected Answer: A
https://cloud.google.com/appengine/docs/standard/php/memcache/using
upvoted 1 times

gaojun 1 year, 8 months ago
Obviously, the answer is A
upvoted 1 times

ehgm 1 year, 11 months ago
Selected Answer: A
Dedicated and shared will resolve the problem, the key is: store all queries in only one key "cached_queries" is not good, we have limits: https://cloud.google.com/appengine/docs/standard/python/memcache
Create a key of each query is better.
upvoted 3 times

vincy2202 2 years ago
A is the correct answer
upvoted 2 times

nqthien041292 2 years ago
Selected Answer: A
Vote A
upvoted 1 times

joe2211 2 years ago
Selected Answer: A
vote A
upvoted 1 times

victory108 2 years, 7 months ago
A. Set the memcache service level to dedicated. Create a key from the hash of the query, and return database values from memcache before issuing a query to Cloud SQL.
upvoted 3 times

un 2 years, 7 months ago
A is correct
upvoted 1 times

Ausias18 2 years, 8 months ago
Answer is A
upvoted 1 times

ga 2 years, 9 months ago
A is correct
upvoted 1 times

BobBui 2 years, 10 months ago
My answer is A
upvoted 1 times

Question #100 *Topic 1*

You need to ensure reliability for your application and operations by supporting reliable task scheduling for compute on GCP. Leveraging Google best practices, what should you do?

A. Using the Cron service provided by App Engine, publish messages directly to a message-processing utility service running on Compute Engine instances.

B. Using the Cron service provided by App Engine, publish messages to a Cloud Pub/Sub topic. Subscribe to that topic using a message-processing utility service running on Compute Engine instances.

C. Using the Cron service provided by Google Kubernetes Engine (GKE), publish messages directly to a message-processing utility service running on Compute Engine instances.

D. Using the Cron service provided by GKE, publish messages to a Cloud Pub/Sub topic. Subscribe to that topic using a message-processing utility service running on Compute Engine instances.

JoeShmoe Highly Voted 4 years, 7 months ago
Answer is B
upvoted 32 times

Smart Highly Voted 4 years, 3 months ago
B is correct. More appropriately: https://cloud.google.com/solutions/reliable-task-scheduling-compute-engine
upvoted 30 times

fraloca 3 years, 5 months ago
https://cloud.google.com/solutions/reliable-task-scheduling-compute-engine#schedule-compute-engine
upvoted 4 times

xaqanik Most Recent 4 months, 2 weeks ago
Selected Answer: B
You can create Cron job using HTTP endpoint, Pub/Sub and App engine.
upvoted 1 times

odacir 7 months ago
Selected Answer: B
Answer is B, but this question is outdated, Today the best practices for cron is Cloud Scheduler: fully managed enterprise-grade cron job scheduler
https://cloud.google.com/scheduler/?gad_source=1&gclsrc=ds&gclsrc=ds
upvoted 15 times

JaimeMS 2 weeks, 2 days ago
Thanks... I was a little confused by this options
upvoted 2 times

JPA210 8 months, 1 week ago
This seems to be an old question, despite B could be the more correct answer, it is not exactly a good one. 'Using the Cron service provided App Engine', the cron service is provided by Cloud Scheduler, not App Engine. App Engine HTTP endpoint can be a target for the cron task.
upvoted 9 times

salim_ 1 year, 1 month ago
Selected Answer: B
https://cloud.google.com/blog/products/gcp/reliable-task-scheduling-on-google-compute-engine
upvoted 3 times

rr4444 1 year, 3 months ago
Something feels missing/broken about this question

Even before comments in discussion that correctly mentioned Cloud Scheduler, which is not mentioned in the question
upvoted 6 times

parthkulkarni998 6 months ago
This is because cloud scheduler is a newly released service which is a replacement to cloud app engine cron service.
upvoted 1 times

dataqueen_3110 1 year, 5 months ago
"By using Cloud Scheduler for scheduling and Pub/Sub for distributed messaging, you can build an application to reliably schedule tasks acr
a fleet of Compute Engine instances." https://cloud.google.com/architecture/reliable-task-scheduling-compute-engine

Answer is B. (Note: It was down to B or D but containerization was not mentioned)
upvoted 2 times

beehive 1 year, 5 months ago
Answer is B.
Cloud Scheduler provides a fully managed, enterprise-grade service that lets you schedule events. After you have scheduled a job, Cloud Scheduler will call the configured event handlers, which can be App Engine services, HTTP endpoints, or Pub/Sub subscriptions.

To run tasks on your Compute Engine instance in response to Cloud Scheduler events, you need to relay the events to those instances. One v
to do this is by calling an HTTP endpoint that runs on your Compute Engine instances. Another option is to pass messages from Cloud Scheduler to your Compute Engine instances using Pub/Sub.
upvoted 8 times

habros 1 year, 6 months ago
Selected Answer: B
A and C are out… messages are to be sent to pub sub and processed using a client. D is overkill for this purpose
upvoted 2 times

megumin 1 year, 7 months ago
Selected Answer: B
B is ok
upvoted 1 times

AzureDP900 1 year, 8 months ago
B is right
upvoted 1 times

Nirca 1 year, 9 months ago
Selected Answer: B
Ans is B https://cloud.google.com/architecture/reliable-task-scheduling-compute-engine
upvoted 1 times

Sbgani 1 year, 9 months ago
Ans is B https://cloud.google.com/architecture/reliable-task-scheduling-compute-engine

refer the examples with diagram
upvoted 1 times

zellck 1 year, 9 months ago
the link points to use Cloud Scheduler, and not Cron service provided by App Engine.
upvoted 3 times

zr79 1 year, 8 months ago
This is the new way to run schedule
upvoted 2 times

FAD04 1 year, 9 months ago
I got this question in exam 01/09/2022
upvoted 5 times

pp0709 1 year, 9 months ago

Selected Answer: D

This solution can be implemented using both A and D
1) With App Engine - https://cloud.google.com/appengine/docs/flexible/nodejs/scheduling-jobs-with-cron-yaml
2) With GKE - https://cloud.google.com/kubernetes-engine/docs/how-to/cronjobs

They ask for best practices and it's well known that GKE (aka containers) is the best practice for building modern infra solution.

Yet another confusing PCA question on the card. Honestly, think the quality of the questions can be mightily improved.

upvoted 1 times

BiddlyBdoyng 1 year, 8 months ago

GKE is too expensive if all you are after is cron scheduling.

upvoted 1 times

pp0709 1 year, 9 months ago

Sorry, can be implemented using both B and D

upvoted 2 times

kapara 10 months, 3 weeks ago

You right, but D is overkill.
So B if the best practices for this task.

upvoted 1 times

medi01 1 year, 1 month ago

A is a bad solution as "send message directly to the utility" is not really reliable, you'd want pub/sub in between.

upvoted 1 times

6721sora 1 year, 9 months ago

B says Appengine.
But Cloud Scheduler is itself a managed service.
To schedule jobs via AppEngine, the cron.yaml has to be used.
It can be done similarly via GKE as well.
This question is confusing

upvoted 2 times

Question #101 *Topic 1*

Your company is building a new architecture to support its data-centric business focus. You are responsible for setting up the network. Your company's mobile and web-facing applications will be deployed on-premises, and all data analysis will be conducted in GCP. The plan is to process and load 7 years of archived .csv files totaling 900 TB of data and then continue loading 10 TB of data daily. You currently have an existing 100-MB internet connection.
What actions will meet your company's needs?

A. Compress and upload both archived files and files uploaded daily using the gsutil λ€"m option.

B. Lease a Transfer Appliance, upload archived files to it, and send it to Google to transfer archived data to Cloud Storage. Establish a connection with Google using a Dedicated Interconnect or Direct Peering connection and use it to upload files daily.

C. Lease a Transfer Appliance, upload archived files to it, and send it to Google to transfer archived data to Cloud Storage. Establish one Cloud VPN Tunnel to VPC networks over the public internet, and compress and upload files daily using the gsutil λ€"m option.

D. Lease a Transfer Appliance, upload archived files to it, and send it to Google to transfer archived data to Cloud Storage. Establish a Cloud VPN Tunnel to VPC networks over the public internet, and compress and upload files daily.

KouShikyou Highly Voted 5 years, 2 months ago

With option A, daily data would take 27 hours.
My answer is B.
How do you think?
upvoted 52 times

xaqanik 10 months, 2 weeks ago
Also with option A you need months to download archive files(900TB)
upvoted 1 times

kumarp6 4 years, 1 month ago
B is correct
upvoted 5 times

Jay_82 4 years ago
ok but dedicated connection is available from 10 GBPS right where as in question it says internet connection is 100 MB, to me D is cor
upvoted 5 times

9xnine 2 years, 6 months ago
Dedicated Interconnect will be a new connection and will not run over the existing internet connection. With dedicated interconnect existing ISP becomes irrelevant. If you were trying to use VPN the existing internet connection would be relevant. Answer is B.
upvoted 5 times

nitinz 3 years, 9 months ago
it is B
upvoted 7 times

malequardos 3 years, 7 months ago
Direct peering is meant only to connect to G Suite Services. Its reference may invalidate the whole answer.
upvoted 2 times

NG123 2 years, 5 months ago
True. B is the most apt answer with just this extra bit "direct peering" raising some confusion.
upvoted 1 times

wk Highly Voted 5 years, 2 months ago
Agree B. 100Mbps connections for 10TB data transfer is takes too long

https://cloud.google.com/solutions/transferring-big-data-sets-to-gcp#close
upvoted 21 times

JJu 5 years ago
not 100Mbps. 100MB
upvoted 3 times

misho 4 years, 6 months ago
even with 100MB internet it's slow. It's 800 Mbps and transfer for 10 TB will take 2 days
upvoted 5 times

bogd 3 years, 10 months ago
There is no such thing as a "100MB" internet connection :) . That must be a speed (per second), and I would guess that the "B" is just a typo (it is highly atypical to measure bandwidth in Bps).
upvoted 5 times

RickMorais Most Recent 5 months, 1 week ago
Selected Answer: B
There is no option than B.
upvoted 1 times

ManojNegi 6 months, 3 weeks ago
Selected Answer: B
B is correct
upvoted 1 times

Devx198912233 10 months, 4 weeks ago

Selected Answer: B

B would be correct as pubsub b service might redeliver messages. When you receive messages in order and the Pub/Sub service redelivers a message with an ordering key, Pub/Sub maintains order by also redelivering the subsequent messages with the same ordering key. The Pub/S service redelivers these messages in the order that it originally received them.

upvoted 1 times

gun123 11 months, 2 weeks ago

Selected Answer: D

B has an option of direct peering too which is not a recommended practise

upvoted 1 times

BisoWafik 1 year ago

Selected Answer: B

B makes sense

upvoted 1 times

odacir 1 year, 1 month ago

Selected Answer: B

B. https://cloud.google.com/architecture/migration-to-google-cloud-transferring-your-large-datasets#time

upvoted 1 times

Arun_m_123 1 year, 2 months ago

Selected Answer: C

I don't think the company needs "dedicated intereconnect" - because it is clearly said that the company wants to do "data-centric" business. Dedicated interconnect is more for having private-access to google cloud.

C seems like a correct option to me

upvoted 1 times

odacir 1 year, 1 month ago

It's B, transferring over 100mbps 10TB daily is not possible, because could take up to 12 days.. You need a better connection.

upvoted 1 times

telp 1 year, 9 months ago

Selected Answer: B

Answer B => Dedicated interconnect will provide a private network with 10gbs. The internet limited to 100 mb is not possible to use cloud VP it will use public internet so be limited for the daily)

upvoted 1 times

sunny2421 1 year, 11 months ago

B is correct.

upvoted 1 times

habros 2 years ago

Selected Answer: B

B. Since it is a new network just sign up for a dedicated line…

upvoted 1 times

megumin 2 years, 1 month ago

Selected Answer: B

B is ok

upvoted 1 times

Balaji_Sakthi 2 years, 1 month ago

its option B. i think

upvoted 1 times

zr79 2 years, 2 months ago

you can not use gsutil to load 10TB daily >>>and then continue loading 10 TB of data daily<<< it will take longer than 24hrs to upload using gsutil

upvoted 2 times

AzureDP900 2 years, 2 months ago

B is the best, VPN doesn't scale very well for huge data

upvoted 2 times

Nirca 2 years, 3 months ago

Selected Answer: B

B is ok

upvoted 1 times

Question #102 *Topic 1*

You are developing a globally scaled frontend for a legacy streaming backend data API. This API expects events in strict chronological order with no repeat data for proper processing.

Which products should you deploy to ensure guaranteed-once FIFO (first-in, first-out) delivery of data?

A. Cloud Pub/Sub alone

B. Cloud Pub/Sub to Cloud Dataflow

C. Cloud Pub/Sub to Stackdriver

D. Cloud Pub/Sub to Cloud SQL

exampanic Highly Voted 4 years, 11 months ago

I believe the answer is B. "Pub/Sub doesn't provide guarantees about the order of message delivery. Strict message ordering can be achieved with buffering, often using Dataflow." https://cloud.google.com/solutions/data-lifecycle-cloud-platform

upvoted 68 times

TiagoM 3 years, 7 months ago

Now Pub/Sub guarantees message order. Until the exam does not change I would pick B.

upvoted 9 times

jask 3 years, 2 months ago

Answer is B. The question is talking about guaranteed-once FIFO delivery of data. Although Pub/sub provides data in order (FIFO) but i does 'at-least' once delivery of data. So, we need Dataflow for deduplication of data.

upvoted 13 times

emirhosseini 2 years, 2 months ago

I believe Pub/Sub now also supports exactly once delivery (in preview):
https://cloud.google.com/pubsub/docs/exactly-once-delivery

upvoted 9 times

melono 2 years, 2 months ago

https://cloud.google.com/pubsub/docs/exactly-once-delivery
reference

upvoted 1 times

melono 2 years, 2 months ago

Pub/Sub supports exactly-once delivery, within a cloud region.
The question states ¨global¨, so needs Dataflow

upvoted 9 times

CosminCiuc 1 year, 10 months ago

I believe that only the frontend is scaled globally. The backend API is the one that requires ordered delivery of the messages and guaranteed-once delivery of data. Currently, Pub/Sub supports ordered delivery within the same region (https://cloud.google.com/pubsub/docs/ordering#receiving_messages_in_order) and exactly-once delivery within the same region (https://cloud.google.com/pubsub/docs/exactly-once-delivery#exactly-once_delivery_guarantees).
The right answer could be A, Pub/Sub alone.

upvoted 3 times

zanfo 2 years, 9 months ago

the correct is B https://cloud.google.com/pubsub/docs/stream-messages-dataflow

upvoted 1 times

xhova Highly Voted 4 years, 7 months ago

B is the answer. CloudSQL is only for storage, to get the messages in order you need timestamp processed in dataflow to arrange them befor putting it in any storage volume. The system described is not querying a db it is expecting a stream of messages only dataflow can correct th order. ACID has no value here because the db is not being queried. You'll not find any documentation on pub/sub order being corrected with a db. See notes below on pub/sub and dataflow using timestamps and windows to ensure order

https://cloud.google.com/pubsub/docs/pubsub-dataflow

upvoted 28 times

selected Most Recent 1 month ago

Selected Answer: B

https://cloud.google.com/pubsub/docs/exactly-once-delivery#regional_considerations

upvoted 1 times

awsgcparch 4 months, 3 weeks ago

Selected Answer: A

Google Cloud Pub/Sub now supports message ordering, which ensures that messages with the same ordering key are delivered in the exact order they were published. This feature addresses the requirement for strict chronological order without the need for additional services.

Key Features of Cloud Pub/Sub with Message Ordering:
Message Ordering: By using ordering keys, Pub/Sub can guarantee that messages are delivered in the order they are published.
Exactly-once Delivery: Pub/Sub supports at-least-once delivery and can be configured to handle duplicate messages.
Scalability and Reliability: Pub/Sub is a fully managed service that scales automatically and ensures high availability.

upvoted 3 times