19040e5 7 months ago

Selected Answer: A

A. Cloud Pub/Sub alone

Cloud Pub/Sub Ordered Delivery: Cloud Pub/Sub natively supports ordered delivery when using the same ordering key. This guarantees that messages with the same key are delivered to subscribers in the order they were published, preventing out-of-order events.
Exactly-once Delivery: Cloud Pub/Sub also offers exactly-once delivery within a region, ensuring that each message is delivered to a subscrib only once.

upvoted 6 times

hitmax87 7 months ago

Selected Answer: B

Pub/Sub provide ordered delivery but doesn't ensure deduplication. Cloud SQL is regional resource and it cant perform custom logic. I go wit

upvoted 1 times

mesodan 9 months, 2 weeks ago

Selected Answer: A

A is correct. While Cloud Dataflow can be used for data processing, it doesn't guarantee FIFO order on its own. Additionally, introducing anot processing layer adds complexity and might not be necessary for this specific requirement.

upvoted 2 times

yas_cloud 10 months, 1 week ago

Selected Answer: B

Dont see a need for SQL here as the answer suggests. Option B is more appropriate

upvoted 1 times

Devx198912233 10 months, 4 weeks ago

Selected Answer: B

B would be correct as pubsub b service might redeliver messages. When you receive messages in order and the Pub/Sub service redelivers a message with an ordering key, Pub/Sub maintains order by also redelivering the subsequent messages with the same ordering key. The Pub/S service redelivers these messages in the order that it originally received them.

upvoted 1 times

Pime13 11 months ago

Selected Answer: B

b, even if pubsub now have exactly-once-delivery this is within a region. question is for a global app.
https://cloud.google.com/pubsub/docs/exactly-once-delivery#exactly-once_delivery_2

upvoted 3 times

theBestStudent 1 year, 1 month ago

Answer is B.
Key words: globally scaled and ensure delivery only once: Pubsub + Dataflow.

If it were only one region, it would be fine to say just pubsub, but it is globally scaled.

upvoted 1 times

Frusci 1 year, 3 months ago

Selected Answer: B

B, you need dataflow to deduplicate. Pub/Sub does "at-least" once delivery.

upvoted 2 times

AL_everyday 1 year, 4 months ago

Selected Answer: A

It should be A

upvoted 1 times

jlambdan 1 year, 6 months ago

Selected Answer: A

https://cloud.google.com/pubsub/docs/exactly-once-delivery

upvoted 2 times

AmarReddy 1 year, 6 months ago
Answer: B
upvoted 1 times

TheCloudGuruu 1 year, 7 months ago
Selected Answer: B
Pub/Sub and Dataflow
upvoted 1 times

Hisayuki 1 year, 8 months ago
Selected Answer: B
Pub/Sub is an at-least-once service. So you need to deduplicate messages with DataFlow as a data pipeline.
upvoted 1 times

Question #103 *Topic 1*

Your company is planning to perform a lift and shift migration of their Linux RHEL 6.5+ virtual machines. The virtual machines are running in an on-premises
VMware environment. You want to migrate them to Compute Engine following Google-recommended practices. What should you do?

A. 1. Define a migration plan based on the list of the applications and their dependencies. 2. Migrate all virtual machines into Compute Engine individually with Migrate for Compute Engine.

B. 1. Perform an assessment of virtual machines running in the current VMware environment. 2. Create images of all disks. Import disks on Compute Engine. 3. Create standard virtual machines where the boot disks are the ones you have imported.

C. 1. Perform an assessment of virtual machines running in the current VMware environment. 2. Define a migration plan, prepare a Migrate for Compute Engine migration RunBook, and execute the migration.

D. 1. Perform an assessment of virtual machines running in the current VMware environment. 2. Install a third-party agent on all selected virtual machines. 3. Migrate all virtual machines into Compute Engine.

kopper2019 Highly Voted 3 years, 5 months ago
Ans) C ,
Migrate for Compute Engine organizes groups of VMs into Waves. After understanding the dependencies of your applications, create runbook that contain groups of VMs and begin your migration!
https://cloud.google.com/migrate/compute-engine/docs/4.5/how-to/migrate-on-premises-to-gcp/overview
upvoted 34 times

technodev Highly Voted 2 years, 11 months ago
Selected Answer: C
I got this question in my exam.
upvoted 15 times

Sur_Nikki 1 year, 7 months ago
Did u passed...? If yes, then Congratulations and let me know the correct answer
upvoted 1 times

3fd692e Most Recent 2 months, 2 weeks ago
Selected Answer: C
Assess, Plan, Migrate. Textbook perfect
upvoted 1 times

salim_ 1 year, 7 months ago
Selected Answer: C
https://cloud.google.com/migrate/compute-engine/docs/4.11/how-to/migrate-on-premises-to-gcp/overview
upvoted 1 times

8d31d36 1 year, 10 months ago

Selected Answer: B

The reason why Option B is preferable over Option C is that it involves creating images of all disks and importing them into Compute Engine, which can significantly reduce the amount of time required for the migration. Additionally, creating standard virtual machines from the importe disks is a straightforward process, and it ensures that the migrated virtual machines are identical to the on-premises virtual machines, which c simplify the migration process and minimize the risk of compatibility issues.

upvoted 2 times

examch 1 year, 11 months ago

Selected Answer: C

C is the correct answer,

Runbooks are created from the Migrate for Compute Engine Manager. The system queries VMware or AWS for VMs and generates a CSV for to edit.

By editing the CSV, you define:

The VMs in a wave.
The order in which those VMs are migrated.
The type and disk space of VMs that are launched on Google Cloud.
Other characteristics that are defined in the Runbook reference.

https://cloud.google.com/migrate/compute-engine/docs/4.8/how-to/organizing-migrations/creating-and-modifying-runbooks#generating_runbook_templates

upvoted 1 times

megumin 2 years, 1 month ago

Selected Answer: C

C is ok

upvoted 1 times

AzureDP900 2 years, 2 months ago

C is most suitable for this use

upvoted 1 times

ACE_ASPIRE 2 years, 4 months ago

I got this question in exam.

upvoted 5 times

AzureDP900 2 years, 5 months ago

C is right, It defines all logical steps to migrate on-premise to google cloud.

upvoted 2 times

meokey 2 years, 8 months ago

Does Ans. C) still valid as of latest GCE 5.0?
in the doc "Migrating VM groups" with version GCE 5.0, I do not see "runbook" anymore which is explained up to version GCE 4.8.
https://cloud.google.com/migrate/compute-engine/docs/5.0/how-to/migrating-vm-groups

upvoted 2 times

gaojun 2 years, 8 months ago

Go for C

upvoted 1 times

[Removed] 2 years, 10 months ago

Selected Answer: C

I got this question on my exam. Answered C.

upvoted 2 times

Sur_Nikki 1 year, 7 months ago

Thanks

upvoted 1 times

haroldbenites 3 years ago
Go for C.
upvoted 1 times

nikiwi 3 years ago
why not A?
seems pretty obvious if you look at the google doc: https://cloud.google.com/migrate/compute-engine/docs/5.0/concepts/lifecycle
upvoted 1 times

atlasga 2 years, 12 months ago
When you are doing cloud migrations, you do migrations in "waves" which are groupings of one or more applications/workloads. Moving machines individually would break things, such as dependencies. This is standard industry practice.
upvoted 3 times

vincy2202 3 years ago
C is the correct answer.
upvoted 1 times

joe2211 3 years ago
Selected Answer: C
vote C

Question #104 *Topic 1*

You need to deploy an application to Google Cloud. The application receives traffic via TCP and reads and writes data to the filesystem. The application does not support horizontal scaling. The application process requires full control over the data on the file system because concurrent access causes corruption. The business is willing to accept a downtime when an incident occurs, but the application must be available 24/7 to support their business operations. You need to design the architecture of this application on Google Cloud. What should you do?

A. Use a managed instance group with instances in multiple zones, use Cloud Filestore, and use an HTTP load balancer in front of the instances.

B. Use a managed instance group with instances in multiple zones, use Cloud Filestore, and use a network load balancer in front of the instances.

C. Use an unmanaged instance group with an active and standby instance in different zones, use a regional persistent disk, and use an HTTP load balancer in front of the instances.

D. Use an unmanaged instance group with an active and standby instance in different zones, use a regional persistent disk, and use a network load balancer in front of the instances.

VishalB Highly Voted 2 years, 11 months ago
Correct Ans : D
Since the Traffic is TCP, Ans A & C gets eliminated as HTTPS load balance is not supported.
B - File storage system is Cloud Firestore which do not give full control, hence eliminated.
D - Unmanaged instance group with network load balance with regional persistent disk for storage gives full control which is required for the migration.
upvoted 62 times

kimharsh 2 years ago
what about the fact that is the unmanaged instance group is not regional , so you can't create it in more than 1 zone ?
upvoted 7 times

Jerryzzyy 10 months, 2 weeks ago
Can we group to running instances in different zones to an unmanaged instance group?
upvoted 1 times

poseidon24 2 years, 10 months ago

almost all good, except for File Storage, is not Cloud Firestore, it is a new service for sharing filesystems across VMs (like a NAS in a traditional infrastructure).

upvoted 10 times

kopper2019 Highly Voted 2 years, 11 months ago

Ans) D , unmanaged instance group as application does not support horizontal scaling and network load balancer as no mention of http traffi

upvoted 28 times

Polosaty Most Recent 3 months ago

In Unmanaged Instance Group instances cannot be in different zones. I think that correct is D but maybe a mistake in the question.

upvoted 1 times

OrangeTiger 5 months ago

Selected Answer: D

Why not a B?
Because the application doesn't support for horizonal scale.
I chose D.

upvoted 1 times

duzapo 9 months, 2 weeks ago

Selected Answer: D

D is corrrect TCP load balancer plus UNMANAGED

upvoted 2 times

TheCloudGuruu 1 year, 1 month ago

Selected Answer: D

must be unmanaged

upvoted 1 times

JC0926 1 year, 3 months ago

Selected Answer: D

Since the application does not support horizontal scaling, a managed instance group is not required. Instead, an unmanaged instance group c
be used to ensure that the application runs on multiple instances in different zones for high availability.
The network load balancer is designed to handle TCP and UDP traffic
The HTTP(S) load balancer is designed specifically for HTTP and HTTPS traffic.

upvoted 15 times

Sur_Nikki 1 year, 1 month ago

Thanks for the apt explanation

upvoted 2 times

omermahgoub 1 year, 6 months ago

An unmanaged instance group allows you to create and manage a group of Compute Engine instances manually, rather than using an autoscaling solution like a managed instance group. This is appropriate for an application that does not support horizontal scaling, as you car manually create and manage the number of instances needed to meet the traffic demands.

To ensure high availability and minimize downtime, you should deploy the instances in different zones and use a regional persistent disk to sto the application's data. This will ensure that the application is still available even if one of the instances or a zone experiences an outage.

A network load balancer should be used in front of the instances to distribute traffic to the instances. A network load balancer is a highly avail and scalable load balancing solution that operates at the network layer and can handle high volumes of traffic. It can also balance traffic acros multiple zones to ensure that the application is always available to users.

upvoted 20 times

omermahgoub 1 year, 6 months ago

Therefore, the correct answer is option D: Use an unmanaged instance group with an active and standby instance in different zones, use a regional persistent disk, and use a network load balancer in front of the instances.

upvoted 10 times

oms_muc 1 year, 6 months ago

Selected Answer: D

Regional Persistent Disk, as App requires full control of filesystem data without concurrent access (block storage vs. file storage (NAS).
https://cloud.google.com/compute/docs/instance-groups
https://cloud.google.com/load-balancing/docs/choosing-load-balancer

upvoted 1 times

surajkrishnamurthy 1 year, 6 months ago

Selected Answer: D

Option A & B eliminated because we cannot use managed instance group since the app does not support Horizontal scaling
Option C > HTTP load balancer is Layer 7 & application is receiving traffic via TCP
Option D > is best answer because we are using Network load balancer Layer 4 which meets the condition "application receives traffic via TC

upvoted 5 times

fiercedog 1 year, 6 months ago

Selected Answer: D

Checking the comparative analysis of storage options, we can see that Filestore is not suitable for the workload, hence A and B are out. C is
because it restricts to HTTP traffic. Answer is D

upvoted 2 times

fiercedog 1 year, 6 months ago

https://cloud.google.com/architecture/storage-advisor#comparative_analysis

upvoted 1 times

megumin 1 year, 7 months ago

Selected Answer: D

D is ok

upvoted 1 times

SerGCP 1 year, 7 months ago

Selected Answer: B

D is not possibile you cannot create a regional unmanaged instance group. https://cloud.google.com/compute/docs/instance-groups/creating
groups-of-unmanaged-instances

upvoted 3 times

KyubiBlaze 1 year, 6 months ago

Downtime is acceptable, Disk is regional, in case of issue the unmanaged instance group can be moved to other zone, disk has data. A&B
not at all an option

upvoted 3 times

minmin2020 1 year, 8 months ago

Selected Answer: D

D. Use an unmanaged instance group with an active and standby instance in different zones, use a regional persistent disk, and use a networ
load balancer in front of the instances.
D is the only option as the application does not support horizontal scaling (no MIG), it needs full control (no filestore) and has TCP traffic (no
HTTP LB).

upvoted 1 times

Santanu_01 1 year, 8 months ago

D is more appropriate solution ---In the ques. it is mentioned done not support horizontal scaling --> hence Unmanaged Instance , and traffic
TCP --- hence N/W
Load balancer

upvoted 1 times

Rajeev26 1 year, 8 months ago

Selected Answer: D

does not support horizontal scaling so MIG not needed. TCP traffic so NW LB is ok

upvoted 1 times

Nirca 1 year, 9 months ago

Selected Answer: D

concurrent access causes corruption - it can not be a managed group

upvoted 1 times

Question #105 *Topic 1*

Your company has an application running on multiple Compute Engine instances. You need to ensure that the application can communicate with an on-premises service that requires high throughput via internal IPs, while minimizing latency. What should you do?

A. Use OpenVPN to configure a VPN tunnel between the on-premises environment and Google Cloud.

B. Configure a direct peering connection between the on-premises environment and Google Cloud.

C. Use Cloud VPN to configure a VPN tunnel between the on-premises environment and Google Cloud.

D. Configure a Cloud Dedicated Interconnect connection between the on-premises environment and Google Cloud.

kopper2019 Highly Voted 3 years, 5 months ago
Ans) D , Reason : high throughput via internal IPs
upvoted 67 times

ShadowLord 2 years, 4 months ago
This is tricky questions , it can be achieved by C and D ... Multiple Computes and Costs .. they are trying to test knowledge on VPN and Tunnels ....
upvoted 2 times

XDevX Highly Voted 3 years, 5 months ago
IMHO the correct answer is D.
Reason: "requires high throughput via internal IPs, while minimizing latency" - both are aspects you cannot guarantee with using VPN traversi the internet.
upvoted 23 times

[Removed] Most Recent 3 months, 3 weeks ago
Selected Answer: D
D seems to be correct
upvoted 1 times

the1dv 11 months, 1 week ago
Selected Answer: D
Should be D as high throughput
upvoted 1 times

Gilbaliano 1 year ago
Selected Answer: D
Justo be D
upvoted 1 times

Gilbaliano 1 year ago
Should be D
upvoted 1 times

kamradamir 1 year ago
Selected Answer: D
It should be D, since need to communicate via Internal IP
upvoted 1 times

thewalker 1 year, 1 month ago

Selected Answer: C

Communication through internal IPs - VPN. So, C

upvoted 1 times

ccpmad 6 months, 1 week ago

yes, you are very smart
but question says high throughput and minimizing latency > Interconnect

upvoted 1 times

[Removed] 11 months, 3 weeks ago

https://cloud.google.com/network-connectivity/docs/interconnect/concepts/overview#:~:text=Also%2C%20Cloud%20Interconnect%20connections%20provide,directly%20cessible%20from%20both%20networks.

upvoted 2 times

duzapo 1 year, 3 months ago

Selected Answer: D

D cause you need high throughput

upvoted 2 times

marcjimz 1 year, 4 months ago

D - high throughput

upvoted 1 times

wooloo 1 year, 4 months ago

How is C marked a correct one?

upvoted 2 times

Danomine416 1 year, 7 months ago

Initially thought 'D' but the question says 'high throughput via internal IPs' so go with VPN answer 'C'

upvoted 2 times

grejao 1 year, 8 months ago

It can´t be VPN, Only interconnect can minimizing latency.

D is the right answer.

upvoted 1 times

zerg0 1 year, 10 months ago

Selected Answer: D

Internal IP + high throughput

upvoted 5 times

WFCheong 1 year, 11 months ago

Selected Answer: D

Ans) D , Reason : high throughput via internal IPs

upvoted 3 times

omermahgoub 1 year, 12 months ago

D. Configure a Cloud Dedicated Interconnect connection between the on-premises environment and Google Cloud.

A Cloud Dedicated Interconnect is a high-bandwidth, low-latency network connection that allows you to connect your on-premises environme to Google Cloud Platform (GCP) using a dedicated network connection. It provides a direct physical connection between your on-premises network and GCP, which can help to reduce latency and increase the throughput of your application.

upvoted 3 times

omermahgoub 1 year, 12 months ago

Option A, using OpenVPN to configure a VPN tunnel between the on-premises environment and Google Cloud, is not a recommended approach. OpenVPN is a free and open-source software application that implements virtual private network (VPN) techniques to create sec point-to-point connections. While OpenVPN can be used to establish a VPN tunnel between an on-premises environment and GCP, it may provide the level of performance and reliability required for high-throughput applications.

Option B, configuring a direct peering connection between the on-premises environment and Google Cloud, is not a recommended approa Direct Peering is a high-bandwidth, low-latency network connection that allows you to connect your on-premises network directly to Goog network. While Direct Peering can be used to connect your on-premises environment to GCP, it is typically used for high-bandwidth worklo such as video streaming and may not be suitable for applications that require low latency.

Option C, using Cloud VPN to configure a VPN tunnel between the on-premises

upvoted 2 times

omermahgoub 1 year, 12 months ago

Cloud VPN, would also involve routing traffic over the public internet and would not provide the low latency and high throughput that yo need.

upvoted 1 times

medi01 1 year, 8 months ago

Yeah, direct peering is high bandwidth and low latency, but may not be suitable for applications that require low latency.

Good job, ChatGPT...

upvoted 2 times

SureshbabuK 2 years ago

Selected Answer: D

Highthoughput connection VPN - No Interconnect - Y
via Internal IP addresses. VPN - yes Interconnect - Y
Interconnect yes for both so D

upvoted 3 times

megumin 2 years, 1 month ago

Selected Answer: D

D is ok

upvoted 1 times

Question #106 *Topic 1*

You are managing an application deployed on Cloud Run for Anthos, and you need to define a strategy for deploying new versions of the application. You want to evaluate the new code with a subset of production traffic to decide whether to proceed with the rollout. What should you do?

A. Deploy a new revision to Cloud Run with the new version. Configure traffic percentage between revisions.

B. Deploy a new service to Cloud Run with the new version. Add a Cloud Load Balancing instance in front of both services.

C. In the Google Cloud Console page for Cloud Run, set up continuous deployment using Cloud Build for the development branch. As part of the Cloud Build trigger, configure the substitution variable TRAFFIC_PERCENTAGE with the percentage of traffic you want directed to a new version.

D. In the Google Cloud Console, configure Traffic Director with a new Service that points to the new version of the application on Cloud Run. Configure Traffic Director to send a small percentage of traffic to the new version of the application.

VishalB Highly Voted 2 years, 11 months ago

 Correct Answer: A
o Each deployment to a service creates a revision. A revision consists of a specific container image, along with environment settings such as environment variables, memory limits, or concurrency value.
o Once the new revision is deployed to a Service you can manage the traffic using MANAGE TRAFFIC option inside the revision tab
 https://cloud.google.com/run/docs/resource-model

upvoted 54 times

omermahgoub Highly Voted 1 year, 6 months ago

The correct answer is A. Deploy a new revision to Cloud Run with the new version. Configure traffic percentage between revisions.

Cloud Run for Anthos allows you to deploy new revisions of your application with a specific percentage of traffic, which allows you to perform gradual rollout of the new version. To do this, you can follow these steps:

Deploy a new revision of your application to Cloud Run with the new version.

In the Cloud Run for Anthos console, navigate to the service that you want to roll out the new version for.

In the "Revisions" tab, you should see the new revision listed alongside the current revision.

Use the traffic percentage slider to specify the percentage of traffic that you want to send to the new revision. You can set the percentage to a small value initially, such as 5%, and gradually increase it over time as you evaluate the new version.

Once you have set the traffic percentage, Cloud Run for Anthos will start directing a portion of the traffic to the new revision, allowing you to evaluate the new version with a subset of production traffic.

upvoted 19 times

omermahgoub 1 year, 6 months ago

Option B, deploying a new service and adding a Cloud Load Balancer instance in front of both services, is not recommended because it would require you to create and manage a separate service for the new version, which would be more complex and less efficient than deploying a new revision.

Option C, using continuous deployment with Cloud Build and substitution variables, is not relevant to this scenario because it involves deploying new versions automatically based on changes to a development branch, rather than manually deploying new revisions with a specific percentage of traffic.

Option D, using Traffic Director, is also not relevant because Traffic Director is used for managing traffic between different services or cluste rather than between revisions of the same service.

upvoted 11 times

Sur_Nikki 1 year, 1 month ago

Extremely convinced by your explanations..Have u given this exam?

upvoted 1 times

baertierchen 1 year ago

Those Answers are surely generated by ChatgGPT

upvoted 1 times

ccpmad 1 week, 3 days ago

Yes, chatgpt did the exam, LOL

upvoted 2 times

Pime13 Most Recent 4 months, 2 weeks ago

Selected Answer: A

https://cloud.google.com/run/docs/rollouts-rollbacks-traffic-migration#gradual

upvoted 2 times

02fc23a 6 months, 3 weeks ago

Selected Answer: A

A:
https://cloud.google.com/run/docs/rollouts-rollbacks-traffic-migration#gradual

upvoted 1 times

odacir 7 months ago

Selected Answer: A

La Opción A - El uso de dashboards predefinidos proporciona una visión inmediata y eficiente del estado del sistema, y la capacidad de agre métricas personalizadas y crear políticas de alertas permite una respuesta rápida y efectiva a los incidentes. Generar un Dashboard por incid introduce complejidad innecesaria.

upvoted 1 times

odacir 7 months ago

Selected Answer: A

https://cloud.google.com/anthos/run/docs/rollouts-rollbacks-traffic-migration

upvoted 1 times

JPA210 8 months, 1 week ago

You all forget that in the exam they expect you to choose the most complete answer, and C is the most efficient and complete one, where you show knowledge of Cloud Build also, and in the real life that is how you should implement this solution.

upvoted 1 times

SSS987 5 months ago

C is saying "from development branch"!

upvoted 2 times

HRS1954 9 months, 2 weeks ago

A is correct - https://cloud.google.com/anthos/run/docs/rollouts-rollbacks-traffic-migration#gradual

upvoted 2 times

BigfootPanda 11 months, 2 weeks ago

Selected Answer: C

C is ok as we want to DEFINE A STRATEGY

upvoted 3 times

TheCloudGuruu 1 year, 1 month ago

Selected Answer: A

Answer is A

upvoted 1 times

zerg0 1 year, 4 months ago

Selected Answer: A

https://cloud.google.com/anthos/run/docs/rollouts-rollbacks-traffic-migration

upvoted 4 times

KM0107 1 year, 6 months ago

Selected Answer: A

Selected A

upvoted 1 times

megumin 1 year, 7 months ago

Selected Answer: A

A is ok

upvoted 1 times

sfsdeniso 1 year, 8 months ago

A is correct
currently there is possibility to use tags to test in production without receiving real traffic: https://cloud.google.com/anthos/run/docs/rollouts-rollbacks-traffic-migration#tags

upvoted 2 times

sfsdeniso 1 year, 8 months ago

whats wrong with C is configuring deployment from 'development branch'
this is supper ugly

upvoted 2 times

zr79 1 year, 8 months ago

you can do a lab on this >>>Deploy Your Website on Cloud Run<<< Manage traffic is there

upvoted 1 times

minmin2020 1 year, 8 months ago

Selected Answer: A

A. Deploy a new revision to Cloud Run with the new version. Configure traffic percentage between revisions.

upvoted 1 times

Question #107 *Topic 1*

You are monitoring Google Kubernetes Engine (GKE) clusters in a Cloud Monitoring workspace. As a Site Reliability Engineer (SRE), you need to triage incidents quickly. What should you do?

A. Navigate the predefined dashboards in the Cloud Monitoring workspace, and then add metrics and create alert policies.

B. Navigate the predefined dashboards in the Cloud Monitoring workspace, create custom metrics, and install alerting software on a Compute Engine instance.

C. Write a shell script that gathers metrics from GKE nodes, publish these metrics to a Pub/Sub topic, export the data to BigQuery, and make a Data Studio dashboard.

D. Create a custom dashboard in the Cloud Monitoring workspace for each incident, and then add metrics and create alert policies.

kopper2019 Highly Voted 2 years, 11 months ago

Ans) A .

upvoted 57 times

DiegoMDZ Highly Voted 2 years, 11 months ago

It's A for me... Create a dashboard for each incident?? I think D isn't a good choice...

upvoted 31 times

bandegg 5 months, 2 weeks ago

Yeah, creating a new dashboard for each incident doesn't seem like the quickest option.

upvoted 3 times

6a8c7ad Most Recent 2 weeks, 5 days ago

quickly would mean custom, and you can call it what you want. D

upvoted 1 times

hitmax87 1 month ago

Selected Answer: D

If you need extended functionality you should create your dashboards, metrics and alerts. Dont mess existing ones.

upvoted 1 times

coolie1234 1 month, 3 weeks ago

Selected Answer: D

Key is custom dashboard

upvoted 1 times

Gino17m 2 months ago

Selected Answer: A

A is correct

upvoted 1 times

mesodan 3 months, 2 weeks ago

Selected Answer: A

A is correct. Option D is highly inefficient and time-consuming. Creating individual dashboards for every incident is impractical and slows dow the triage process.

upvoted 1 times

OrangeTiger 5 months ago

I will go with A.
uhhm, opinions are divided. In such cases, Q is often not good.

upvoted 1 times

hzaoui 5 months ago

Selected Answer: A

Explanation: Cloud Monitoring provides predefined dashboards for monitoring GKE clusters, which facilitate an immediate and comprehensiv view of cluster performance and health. As an SRE, utilizing these dashboards helps triage incidents quickly. You can also add additional met that are pertinent to the incident and create alert policies that will notify you when specific conditions indicative of an incident are met. This strategy allows for the proactive monitoring of incidents and rapid response when necessary.

upvoted 3 times

SSS987 5 months ago

Ans: D. Although creating dashboard per incident sounds confusing and inefficient, it is still better than the impossible option A as we can't ec or add metrics to a predefined dashboard. Inefficient option vs Impossible Option - Inefficient one is ok!

upvoted 3 times

Gino17m 2 months ago

Optipn A is possible. You can't add widgets to predefined dashboard but you can create alert policies based on metrics. Although the structure of the question is actually misleading

upvoted 2 times

ade7cae 6 months, 1 week ago

Selected Answer: A

A is correct. For D, creating dashboards for each incident isn't practical

upvoted 1 times

spuyol 6 months, 1 week ago
Currently you can not modify pre-defined dashboards.
Has no sense to create a dashboard for each incident.
Conclusion: no answer is good enough
upvoted 1 times

brentc 7 months, 1 week ago
Selected Answer: D
We do have dashboards for each incident
upvoted 2 times

TopTalk 8 months, 3 weeks ago
Selected Answer: D
"You can't delete or modify the automatically-created dashboards; however, when support for copying the dashboard exists, you can modify t
copy. In general, you can also copy charts on a predefined dashboard to a dashboard that you create. Dashboards that you create are custon
dashboards. Custom dashboards let you display information that is of interest to you, organized in a way that's useful to you. "
https://cloud.google.com/monitoring/charts/predefined-dashboards
upvoted 3 times

ArtistS 7 months ago
Create a dashboard for each inc ????? serious
upvoted 1 times

TopTalk 9 months, 1 week ago
Selected Answer: D
"To view the chart associated with an alerting policy and information about incidents in the same context as your metric data, add alert charts
and incident widgets to your CUSTOM dashboard." https://cloud.google.com/monitoring/dashboards/alerts-and-incidents
upvoted 3 times

duzapo 9 months, 2 weeks ago
Selected Answer: A
Ans A any sense a dashboar per incident
upvoted 1 times

heretolearnazure 9 months, 3 weeks ago
A is correct
upvoted 1 times

Question #108 *Topic 1*

You are implementing a single Cloud SQL MySQL second-generation database that contains business-critical transaction data. You want to ensure that the minimum amount of data is lost in case of catastrophic failure. Which two features should you implement? (Choose two.)

A. Sharding

B. Read replicas

C. Binary logging

D. Automated backups

E. Semisynchronous replication

kopper2019 Highly Voted 2 years, 11 months ago
Ans) C and D
Cloud SQL. If you use Cloud SQL, the fully managed Google Cloud MySQL database, you should enable automated backups and binary logg
for your Cloud SQL instances. This allows you to perform a point-in-time recovery, which restores your database from a backup and recovers
to a fresh Cloud SQL instance
upvoted 37 times

HenkH 1 year, 7 months ago

And: a read-replica won't help against "catastrophic failures" like accidental deletions

upvoted 5 times

RVivek 1 year, 4 months ago

catastrophic failure means disaster like a zonal datacenter level failure or regional failure

upvoted 4 times

victory108 Highly Voted 2 years, 11 months ago

C. Binary logging
D. Automated backups

upvoted 11 times

[Removed] Most Recent 5 months, 3 weeks ago

CD

Binary Logging: Binary logging in MySQL records changes to the database. It can be used for backup and replication, and it's essential for pc in-time recovery. With binary logging, you can roll your database forward to any point in time, minimizing data loss.

Automated Backups: Automated backups periodically take a snapshot of your database. In the event of a catastrophic failure, you can restore your database to the state it was in at the time of the last backup. This can also help minimize data loss.

While read replicas and semisynchronous replication can enhance availability and performance, they do not directly minimize data loss. Also, you cannot create a read replica without enabling Automated backups and Enable binary logging

Sharding can improve performance but it's not directly aimed at data loss prevention.

https://cloud.google.com/sql/docs/mysql/backup-recovery/backups

upvoted 3 times

odacir 7 months ago

Selected Answer: CD

Prerequisites for creating a read replica

Before you can create a read replica of a primary Cloud SQL instance, the instance must meet the following requirements:

Automated backups must be enabled.
Binary logging must be enabled which requires point-in-time recovery to be enabled. Learn more about the impact of these logs.
At least one backup must have been created after binary logging was enabled.
https://cloud.google.com/sql/docs/mysql/replication#requirements

upvoted 5 times

someone2011 9 months, 1 week ago

CD
Before being able to create a read replica, you have to make sure "binary logging and automated backup" are enabled. So picking only D or C without the other one makes no sense.
https://cloud.google.com/sql/docs/mysql/replication/create-replica

upvoted 1 times

HRS1954 9 months, 2 weeks ago

The correct answers are C. & D.

Binary logging is a feature of MySQL that records all changes made to the database. This log can be used to restore the database to a previou state in case of a failure.

Automated backups are regularly scheduled backups of the database. They are the most reliable way to ensure that data is not lost in case of catastrophic failure.

upvoted 4 times

heretolearnazure 9 months, 3 weeks ago

C and D. Read replicas wont work in this case.

upvoted 1 times

FaizAhmed 11 months, 3 weeks ago

C. Binary logging
D. Automated backups

upvoted 1 times

taer 1 year, 2 months ago

Selected Answer: CD

Binary logging records changes to the data, which can help you recover data and minimize data loss during an unexpected failure. Automatec backups create regular backups of your database, allowing you to restore the database to a specific point in time in case of a catastrophic failure.

upvoted 3 times

abbottWang 1 year, 3 months ago

Selected Answer: CD

backup data automatically

upvoted 1 times

telp 1 year, 3 months ago

CD => the answer B is for performance issue. The question focus on data loss prevention.

upvoted 1 times

medi01 1 year, 1 month ago

So you are going with a SINGLE instance of MySQL for a critical business application.

upvoted 1 times

okixavi 1 year, 4 months ago

Selected Answer: BD

B and D:
No need to explain D, but B... here is why
When you set up a read replica, automaticaly binary logging is activated. Then, in case of desaster, you can promote manually a read replica a it will have all data before the desaster occurs.

upvoted 3 times

r1ck 1 year, 4 months ago

sure, binary logging starts Automatically upon configuring read-replica??
- Don't think so,
https://cloud.google.com/sql/docs/mysql/replication/create-replica

upvoted 2 times

Jeena345 1 year, 4 months ago

B and D are correct answers as per below reference,
1. Before you can create a read replica of a primary Cloud SQL instance, the instance must meet the following requirements:

Automated backups must be enabled.

2. Binary logging must be enabled which requires point-in-time recovery to be enabled. Learn more about the impact of these logs.

3. At least one backup must have been created after binary logging was enabled.

It means creating read replica already covers binary logging.

Please read the following references for more information

https://cloud.google.com/solutions/cloud-sql-mysql-disaster-recovery-complete-failover-fallback

https://medium.com/google-cloud/cloud-sql-recovering-from-regional-failure-in-10-minutes-or-less-mysql-fc055540a8f0

Replication in Cloud SQL | Cloud SQL for MySQL | Google Cloud

upvoted 1 times

mmathiou 1 year, 1 month ago

Yes, you are correct that creating a read replica requires binary logging to be enabled on the primary instance. However, the purpose of a r replica is to scale read traffic and offload it from the primary instance, not to prevent data loss in case of catastrophic failure. While enablin binary logging is a requirement for creating a read replica, it is not the primary purpose of a read replica. IMO the two features that should l implemented to ensure minimum data loss in case of catastrophic failure are Binary logging and Automated backups.

upvoted 3 times

examch 1 year, 5 months ago

Selected Answer: CD

C and D are correct answers,

Backups help you restore lost data to your Cloud SQL instance. Additionally, if an instance is having a problem, you can restore it to a previou state by using the backup to overwrite it. Enable automated backups for any instance that contains necessary data. Backups protect your dat from loss or damage.

Enabling automated backups, along with binary logging, is also required for some operations, such as clone and replica creation.

https://cloud.google.com/sql/docs/mysql/backup-recovery/backups#what_backups_provide

upvoted 2 times

omermahgoub 1 year, 6 months ago

C. Binary logging
D. Automated backups

Binary logging is a feature of MySQL that records all changes made to the database in a binary log file. By enabling binary logging on your Cl SQL instance, you can use the log file to recover your database in case of catastrophic failure.

Automated backups are a feature of Cloud SQL that allows you to automatically create and retain backups of your database. By enabling automated backups, you can restore your database in case of catastrophic failure or other data loss events.

Option A, sharding, is not a recommended approach. Sharding is a technique for distributing data across multiple servers to improve

Question #109 *Topic 1*

You are working at a sports association whose members range in age from 8 to 30. The association collects a large amount of health data, such as sustained injuries. You are storing this data in BigQuery. Current legislation requires you to delete such information upon request of the subject. You want to design a solution that can accommodate such a request. What should you do?

A. Use a unique identifier for each individual. Upon a deletion request, delete all rows from BigQuery with this identifier.

B. When ingesting new data in BigQuery, run the data through the Data Loss Prevention (DLP) API to identify any personal information. As part of the DLP scan, save the result to Data Catalog. Upon a deletion request, query Data Catalog to find the column with personal information.

C. Create a BigQuery view over the table that contains all data. Upon a deletion request, exclude the rows that affect the subject's data from this view. Use this view instead of the source table for all analysis tasks.

D. Use a unique identifier for each individual. Upon a deletion request, overwrite the column with the unique identifier with a salted SHA256 of its value.

milan74 Highly Voted 3 years, 5 months ago

According to me, the question states "The association collects a large amount of health data, such as sustained injuries." and the nuance on t word such => " Current legislation requires you to delete "SUCH" information upon request of the subject. " So from that point of view the question is not to delete the entire user records but specific data related to personal health data. With DLP you can use InfoTypes and InfoTyp detectors to specifically scan for those entries and how to act upon them (link https://cloud.google.com/dlp/docs/concepts-infotypes)
I would say B.

upvoted 84 times

AmitAr 2 years, 7 months ago

(A) - Primary task is "legislation requires you to delete" .. and B is not deleting.
only A is deleting

upvoted 11 times

BeCalm 1 year, 9 months ago

Deletion is implied in "Upon a deletion request, query Data Catalog to find the column with personal information."

upvoted 1 times

zanfo 2 years, 9 months ago

I want to delete all the informations about the user, not only those individuate by DLP. ALL THE INFORMATIONS of the users...B is not corr the correct is A

upvoted 8 times

Ishu_awsguy 2 years, 4 months ago

There is no need of DLP.
All the data is sensitive but only upon user request it needs deletion.
So A should be the correct answer.

upvoted 12 times

Arad 3 years ago

as PhilipKoku mentioned below:
A) is the correct answer. B) is only masking the data and then when a request is received, it identified the record but it doesn't delete it. D) masking the ID.

upvoted 12 times

mgm7 3 years ago

B is not masking the data but identifying where it is to take action on at later date if required

upvoted 6 times

XDevX Highly Voted 3 years, 5 months ago

IMHO a) is the correct answer because it is easier to operate. The question is not how to mask data and so on but just to delete data on requ so I don't think that we have to use for just the deletion of specific data DLP.

upvoted 35 times

andyk87 Most Recent 2 days, 3 hours ago

Selected Answer: B

Option B is better when the requirement is to delete only the PII health data, not all data related to the individual.

upvoted 1 times

Sephethus 6 months ago

It had better be A, if not then you're not a good organization

upvoted 2 times

hitmax87 7 months ago

Selected Answer: B

Data Loss Prevention must have!

upvoted 1 times

Gino17m 7 months, 2 weeks ago

Selected Answer: B

I vote for B.
I had some doubts whether A was correct, but:
- I'm not convinced by the argument "only A talks about deleting" (it would be too easy if it was about choosing an answer containing the wor "delete" ;)
- the question says "design a solution that can accommodate such a request" - I'm not very fluent in english, but "accommodate" imho mean more "facilitate" than "accomplish" here
- I think that the task is about deleting health data not everything related with unique identifier
- Data Catalog allows you to manage data, knowing in which datasets and in which tables what data is stored. Answer "A" somehow imposes the data model - each table with data related to a given individual must contain the ID of this individual (in a real data model this does not hav be the case).

upvoted 3 times

Djenko 9 months ago

Selected Answer: A

Should be A)

upvoted 2 times

mesodan 9 months, 2 weeks ago

Selected Answer: A

A is correct. As for option B: While DLP is valuable for identifying sensitive data, it might not be sufficient for this specific case. DLP cannot necessarily determine an individual's right to deletion based solely on data classification. Additionally, relying on Data Catalog to store the res adds unnecessary complexity and potential inconsistencies.

upvoted 3 times

Gall 10 months, 2 weeks ago

Selected Answer: B

B. The A removes all data, not SUCH only.

upvoted 2 times

NoCrapEva 10 months, 3 weeks ago

Selected Answer: A

Ans. B assumes you will delete the Personal Information found in the Catalog... Some people are reading GDPR into this question (we are not told what country and what legislation). The question states you must delete all information (not just personal informarion) on request. Ans B is red herring !
Answer must = A

upvoted 2 times

Roro_Brother 11 months, 3 weeks ago

Selected Answer: A

(A) - Primary task is "legislation requires you to delete" .. and B is not deleting.
only A is deleting

upvoted 5 times

Jconnor 1 year ago

Well, A would delete all rows with the identifier, I guess including the ones that are not confidential, also what does it mean unique identifier? each user is unique already. Ridiculous. B would identify the columns that contain personal data, but B is prone to errors as changes in legisla of what is consider injury would be excluded and all data would need to be re- ingested. Unfortunately B is closer and less damaging than A.

upvoted 1 times

thewalker 1 year, 1 month ago

Selected Answer: B

Either A or B is the answer.
A - will delete all the info about the subject, which is not the intension. Only the sensitive data to be deleted. Hence, B.

upvoted 2 times

Anubhav451 1 year, 1 month ago

B is correct.. Check in chatGPT also.

upvoted 1 times

DA95 1 year ago

chatgpt select A
The most appropriate solution for accommodating the deletion request of personal health data stored in BigQuery, as per current legislatio would be:

A. Use a unique identifier for each individual. Upon a deletion request, delete all rows from BigQuery with this identifier.

Here's why this approach is suitable:

Unique Identifier: Assigning a unique identifier to each individual is a standard practice in managing and querying datasets. It helps in precisely identifying and isolating records associated with a specific individual.

Direct Deletion of Rows: Upon receiving a deletion request, you can directly delete all rows associated with the individual's unique identifie This approach ensures that the data is completely removed from your dataset, complying with the legislative requirement to delete persona information upon request.

upvoted 1 times

JPA210 1 year, 2 months ago

From one side A is an easy way, low effort, to implement this solution, but if we think like an architect and like an exam question, B is more complete and a better solution, since it can mask all the sensitive information, not only for the users that request it, but for all, which is a best practice.

upvoted 1 times

AdityaGupta 1 year, 2 months ago

Selected Answer: B

We do not need to delete entire recrod of sports person but some health information collected by association. B would be correct answer.

upvoted 2 times

TopTalk 1 year, 3 months ago

Selected Answer: B

The problem I see with A is that it doesn't offer you a way to find the original subject's information once they request for their information to b deleted (no mapping from the unique identifier back to their person). Only B offers the solution design for this ability. The deletion step may nc be included in B, but the ability to delete is always present. You're designing the ability to accommodate the request, which is to look up the individual who is asking for their information to be deleted.

upvoted 1 times